

Bitcoin Mechanics

Yu Feng
University of California, Santa Barbara



Slides adapted from Stanford CS251

Digital Signatures

Physical signatures: bind transaction to author

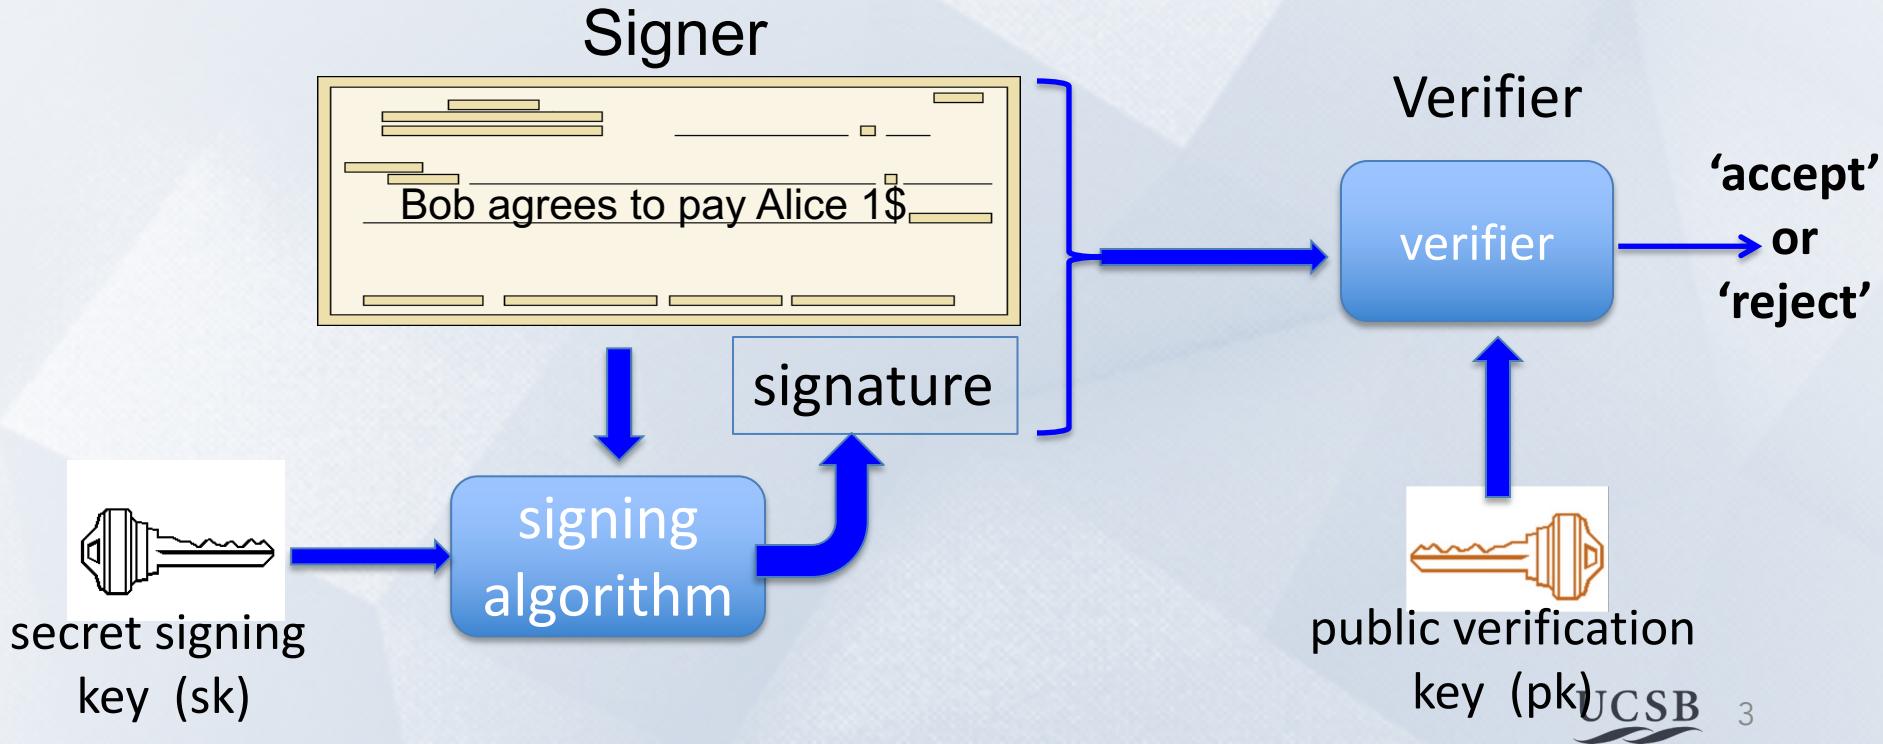


Problem in the digital world:

anyone can copy Bob's signature from one doc to another

Digital signatures

Solution: make signature depend on document



Digital signatures: syntax

Def: a signature scheme is a triple of algorithms:

- $\text{Gen}()$: outputs a key pair (pk, sk)
- $\text{Sign}(\text{sk}, \text{msg})$ outputs sig. σ
- $\text{Verify}(\text{pk}, \text{msg}, \sigma)$ outputs 'accept' or 'reject'

Secure signatures: (informal)

Adversary who sees signatures **on many messages** of his choice,
cannot forge a signature on a new message.

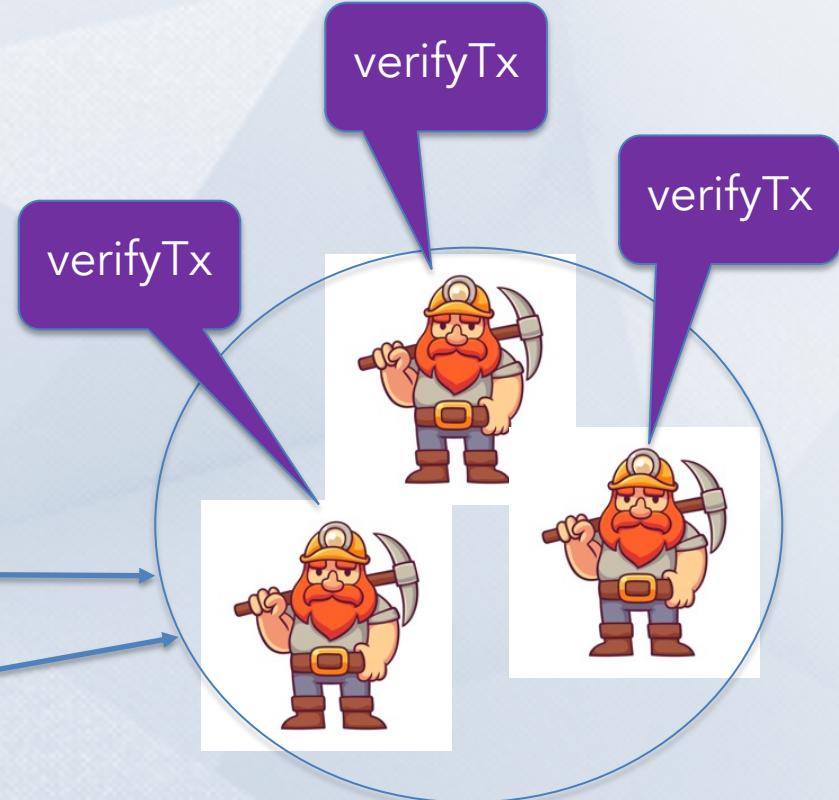
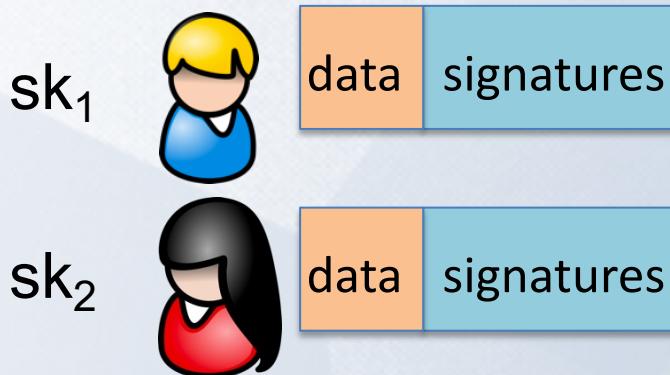
Families of signature schemes

1. RSA signatures (old ... not used in blockchains):
 - long sigs and public keys (≥ 256 bytes), fast to verify
2. Discrete-log signatures: Schnorr and ECDSA (Bitcoin, Ethereum)
 - short sigs (48 or 64 bytes) and public key (32 bytes)
3. BLS signatures: 48 bytes, aggregatable, easy threshold
(Ethereum 2.0, Chia, Dfinity)
4. Post-quantum signatures: long (≥ 600 bytes)

Signatures on the blockchain

Signatures are used everywhere:

- ensure Tx authorization,
- governance votes,
- consensus protocol votes.



In summary ...

Digital signatures: (Gen, Sign, Verify)

$\text{Gen}() \rightarrow (\text{pk}, \text{sk}),$

$\text{Sign}(\text{sk}, m) \rightarrow \sigma,$ $\text{Verify}(\text{pk}, m, \sigma) \rightarrow \text{accept/reject}$

signing key

verification key

Cryptography Background

(1) cryptographic hash functions

An efficiently computable function $H: M \rightarrow T$
where $|M| \gg |T|$



Collision resistance

Def: a **collision** for $H: M \rightarrow T$ is pair $x \neq y \in M$ s.t. $H(x) = H(y)$

$|M| \gg |T|$ implies that many collisions exist

Def: a function $H: M \rightarrow T$ is **collision resistant** if it is “hard” to find even a single collision for H (we say H is a CRF)

Example: **SHA256**: $\{x : \text{len}(x) < 2^{64} \text{ bytes}\} \rightarrow \{0, 1\}^{256}$

(output is 32 bytes)

Application: committing to data on a blockchain

Alice has a large file m . She posts $h = H(m)$ (32 bytes)

Bob reads h . Later he learns m' s.t. $H(m') = h$

H is a CRF \Rightarrow Bob is convinced that $m' = m$
(otherwise, m and m' are a collision for H)

We say that $h = H(m)$ is a **binding commitment** to m

(note: not hiding, h may leak information about m)

Committing to a list (of transactions)

Alice has $S = (m_1, m_2, \dots, m_n)$

32 bytes

Goal:

- Alice posts a short binding commitment to S , $h = \text{commit}(S)$
- Bob reads h . Given $(m_i, \text{proof } \pi_i)$ can check that $S[i] = m_i$
Bob runs $\text{verify}(h, i, m_i, \pi_i) \rightarrow \text{accept/reject}$

security: adv. cannot find (S, i, m, π) s.t. $m \neq S[i]$ and
 $\text{verify}(h, i, m, \pi) = \text{accept}$ where $h = \text{commit}(S)$

Merkle tree

(Merkle 1989)

commitment



h

Merkle tree
commitment

$m_1 \ m_2 \ m_3 \ m_4 \ m_5 \ m_6 \ m_7 \ m_8$



list of values S

Goal:

- commit to list S of size n
- Later prove $S[i] = m_i$

Merkle tree

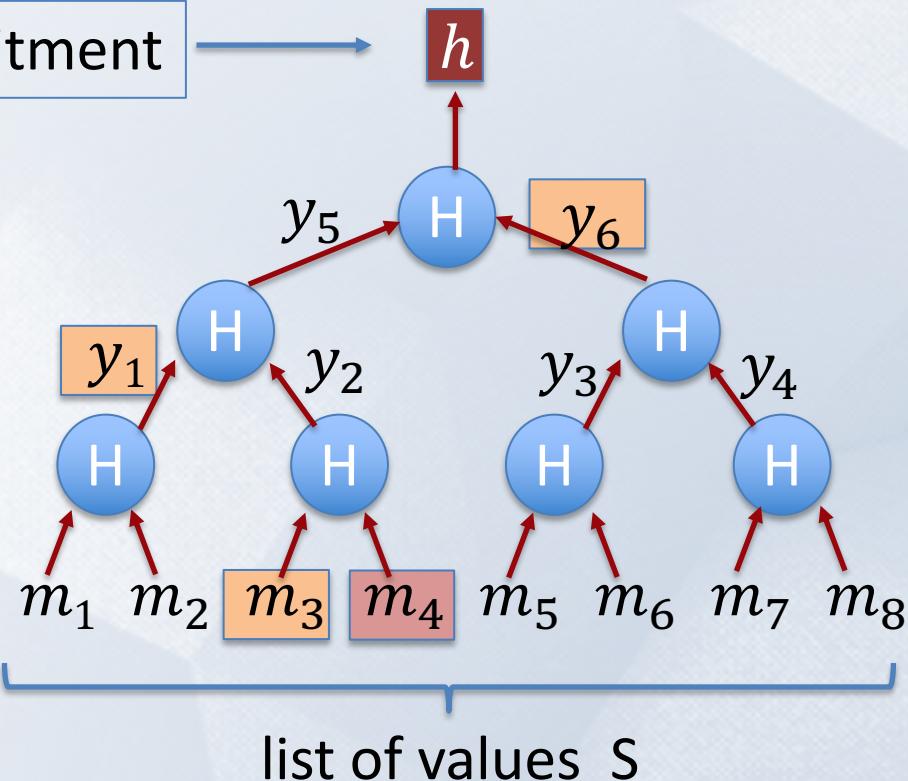
(Merkle 1989)

[simplified]

commitment



h



Goal:

- commit to list S of size n
- Later prove $S[i] = m_i$

To prove $S[4] = m_4$,
proof $\pi = (m_3, y_1, y_6)$

length of proof: $\log_2 n$

Merkle tree

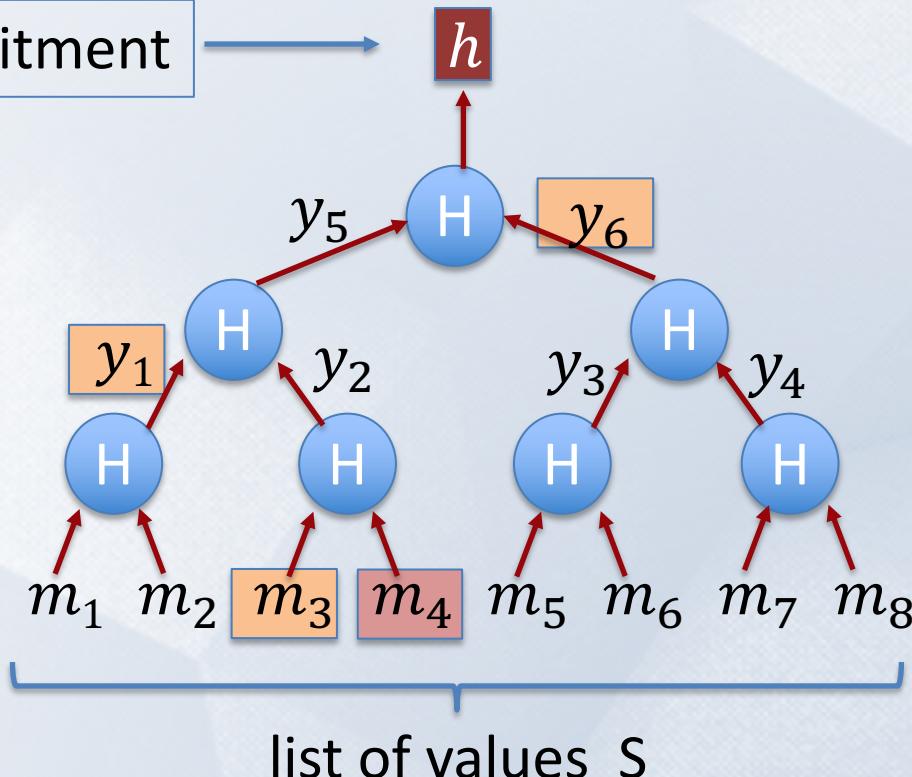
(Merkle 1989)

[simplified]

commitment



h



To prove $S[4] = m_4$,
proof $\pi = (m_3, y_1, y_6)$

Bob does:

$y_2 \leftarrow H(m_3, m_4)$
 $y_5 \leftarrow H(y_1, y_2)$
 $h' \leftarrow H(y_5, y_6)$
accept if $h = h'$

Merkle tree (Merkle 1989)

Thm: For a given n : if H is a CRF then
adv. cannot find (S, i, m, π) s.t. $|S| = n$, $m \neq S[i]$,
 $h = \text{commit}(S)$, and $\text{verify}(h, i, m, \pi) = \text{accept}$

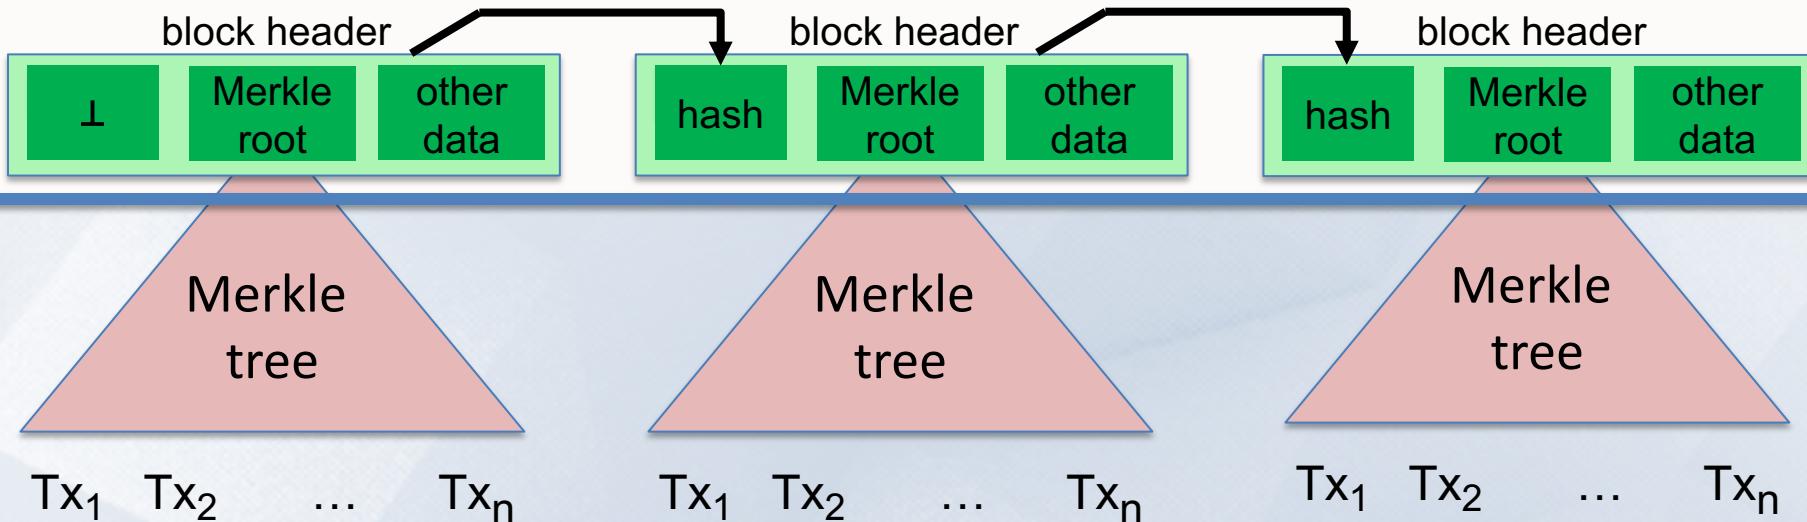
(to prove, prove the contra-positive)

How is this useful? To post a block of transactions S on chain suffices to only write $\text{commit}(S)$ to chain. Keeps chain small.

⇒ Later, can prove contents of every Tx.

Abstract block chain

blockchain



Merkle proofs are used to prove that a Tx is “on the block chain”

Another application: proof of work

Goal: computational problem that

- takes time $\Omega(D)$ to solve, but
- solution takes time $O(1)$ to verify

(D is called the **difficulty**)

How? $H: X \times Y \rightarrow \{0, 1, 2, \dots, 2^n - 1\}$ e.g. $n = 256$

- puzzle: input $x \in X$, output $y \in Y$ s.t. $H(x, y) < 2^n/D$
- verify(x, y): accept if $H(x, y) < 2^n/D$

Another application: proof of work

Thm: if H is a “random function” then the best algorithm requires D evaluations of H in expectation.

Note: this is a parallel algorithm

⇒ the more machines I have, the faster I solve the puzzle.

Proof of work is used in some consensus protocols (e.g., Bitcoin)

Bitcoin uses $H(x, y) = \text{SHA256}(\text{SHA256}(x \cdot y))$

Bitcoin mechanics



This lecture: Bitcoin mechanics

Total market value:

Oct. 2008: paper by Satoshi Nakamoto
Jan. 2009: Bitcoin network launched



This lecture: Bitcoin mechanics

user facing tools (cloud servers)

applications (DAPPs, smart contracts)

Execution engine (blockchain computer)

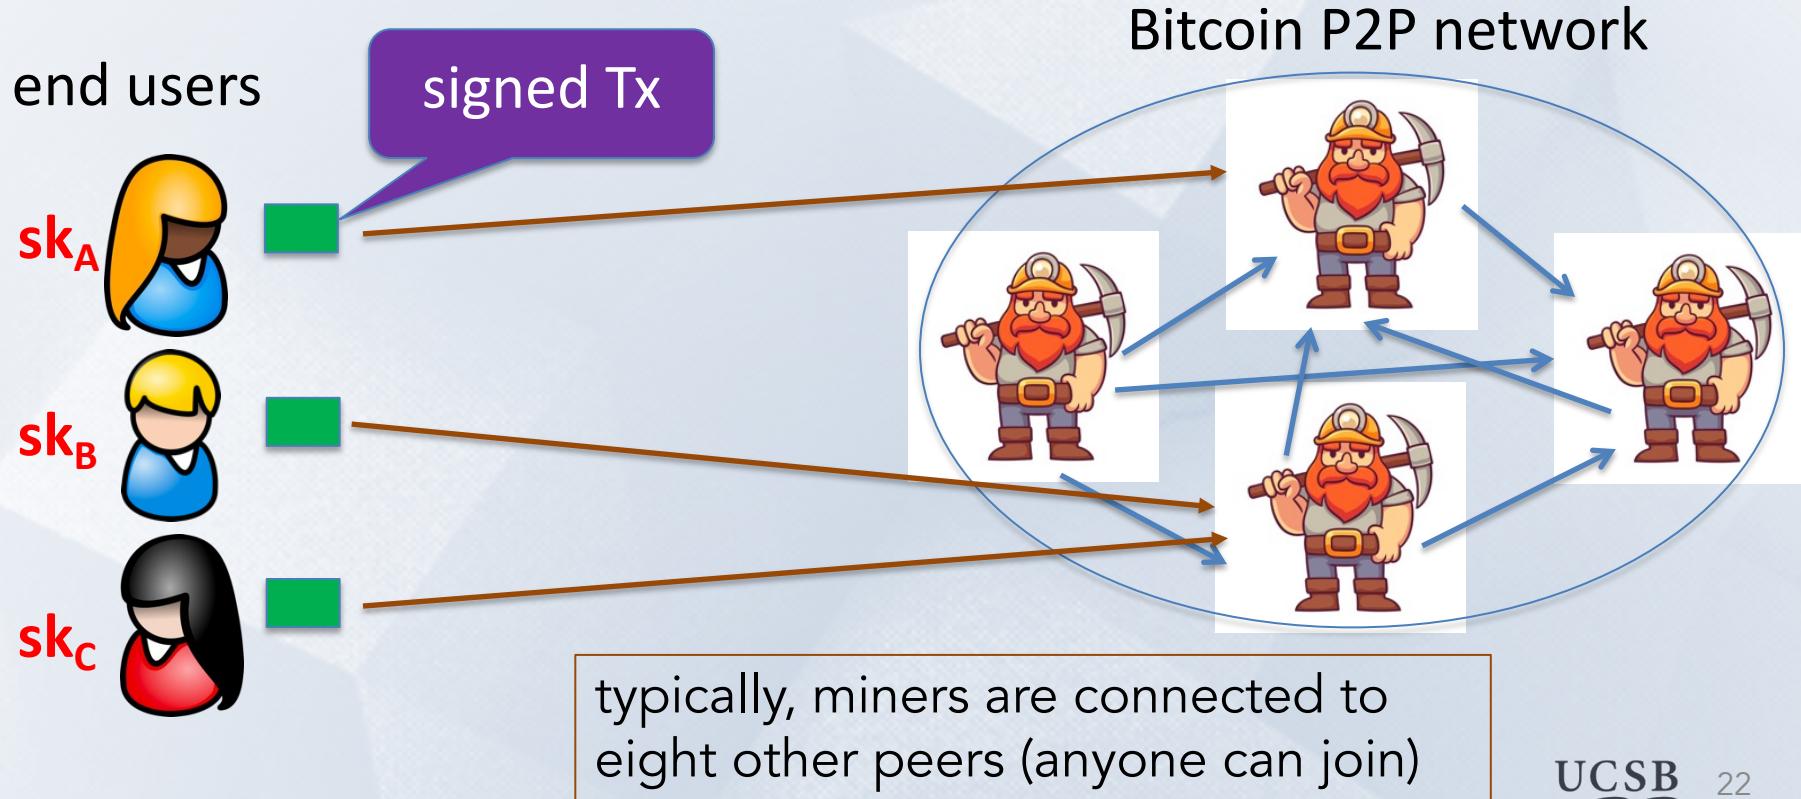
today

Sequencer: orders transactions

Data Availability / Consensus Layer

next week

First: overview of the Bitcoin consensus layer

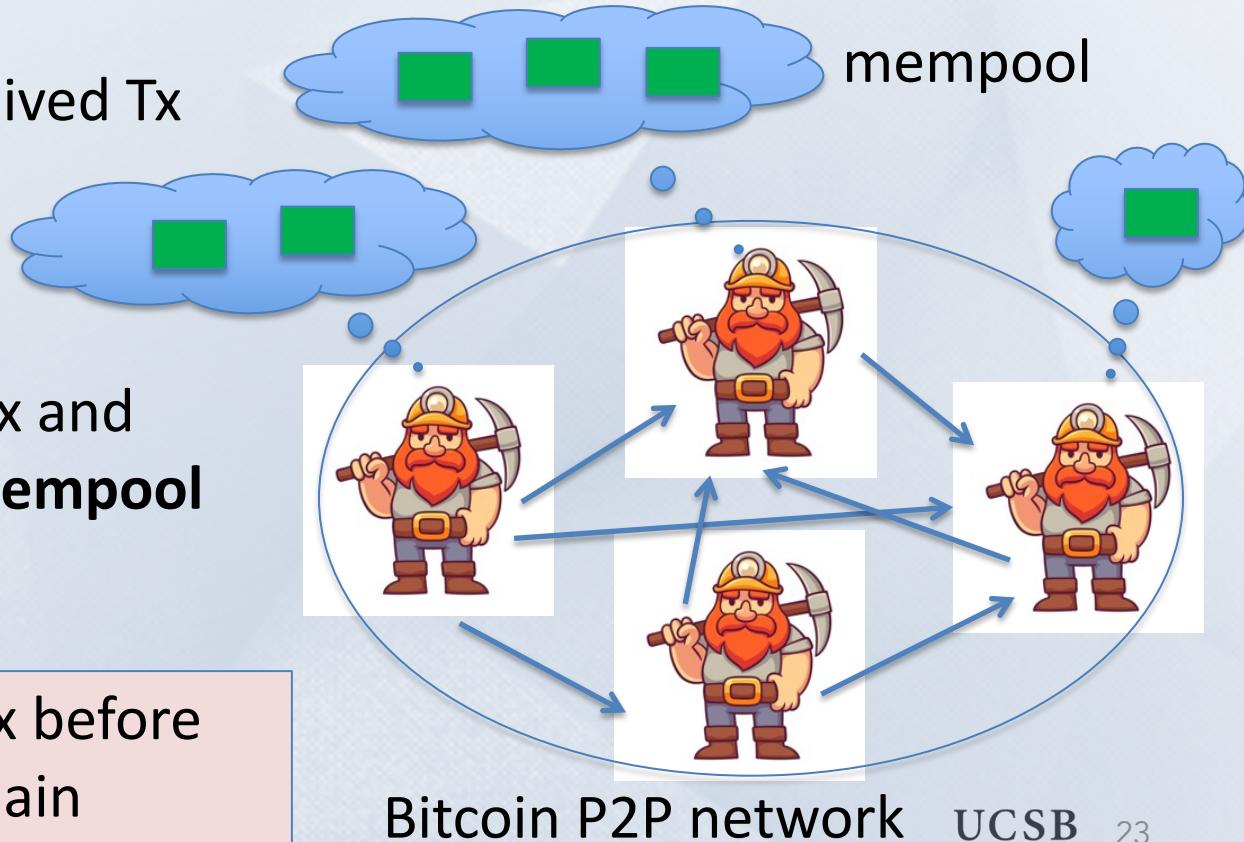


First: overview of the Bitcoin consensus layer

miners broadcast received Tx
to the P2P network

every miner:
validates received Tx and
stores them in its **mempool**
(unconfirmed Tx)

note: miners see all Tx before
they are posted on chain



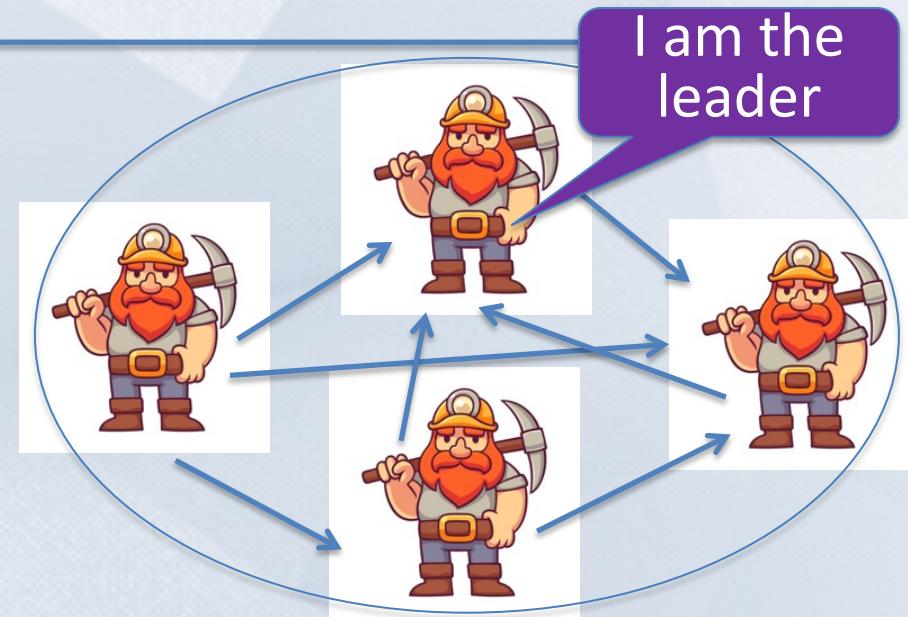
First: overview of the Bitcoin consensus layer

blockchain



Every ≈10 minutes:

- Each miner creates a candidate block from Tx in its mempool
- a “random” miner is selected (how: next week), and broadcasts its block to P2P network
- all miners validate new block



Bitcoin P2P network

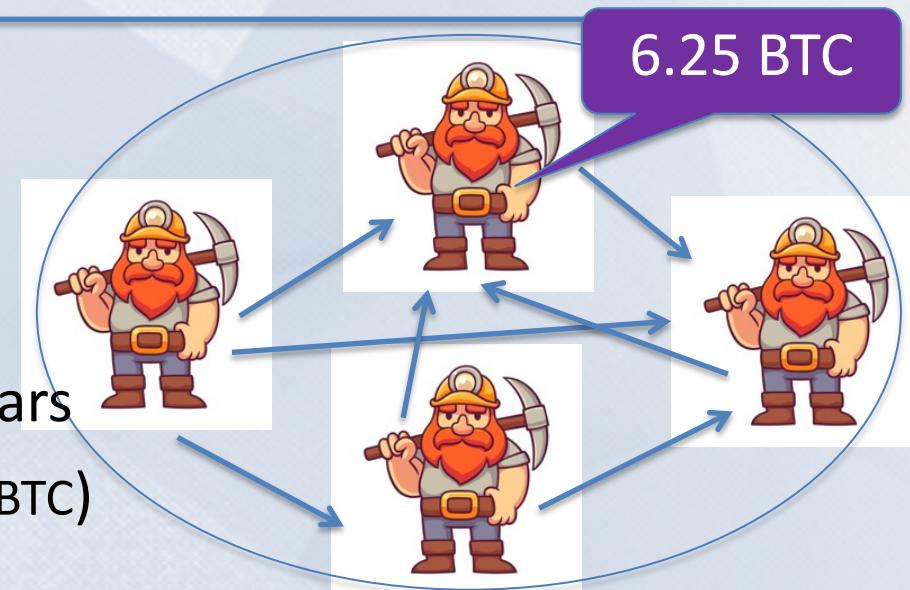
First: overview of the Bitcoin consensus layer

blockchain



Selected miner is paid 6.25 BTC
in **coinbase Tx** (first Tx in the block)

- only way new BTC is created
- block reward halves every four years
⇒ max 21M BTC (currently 19.6M BTC)



note: miner chooses order of Tx in block

Properties (very informal)

Next week:

Safety / Persistence:

- to remove a block, need to convince 51% of mining power *

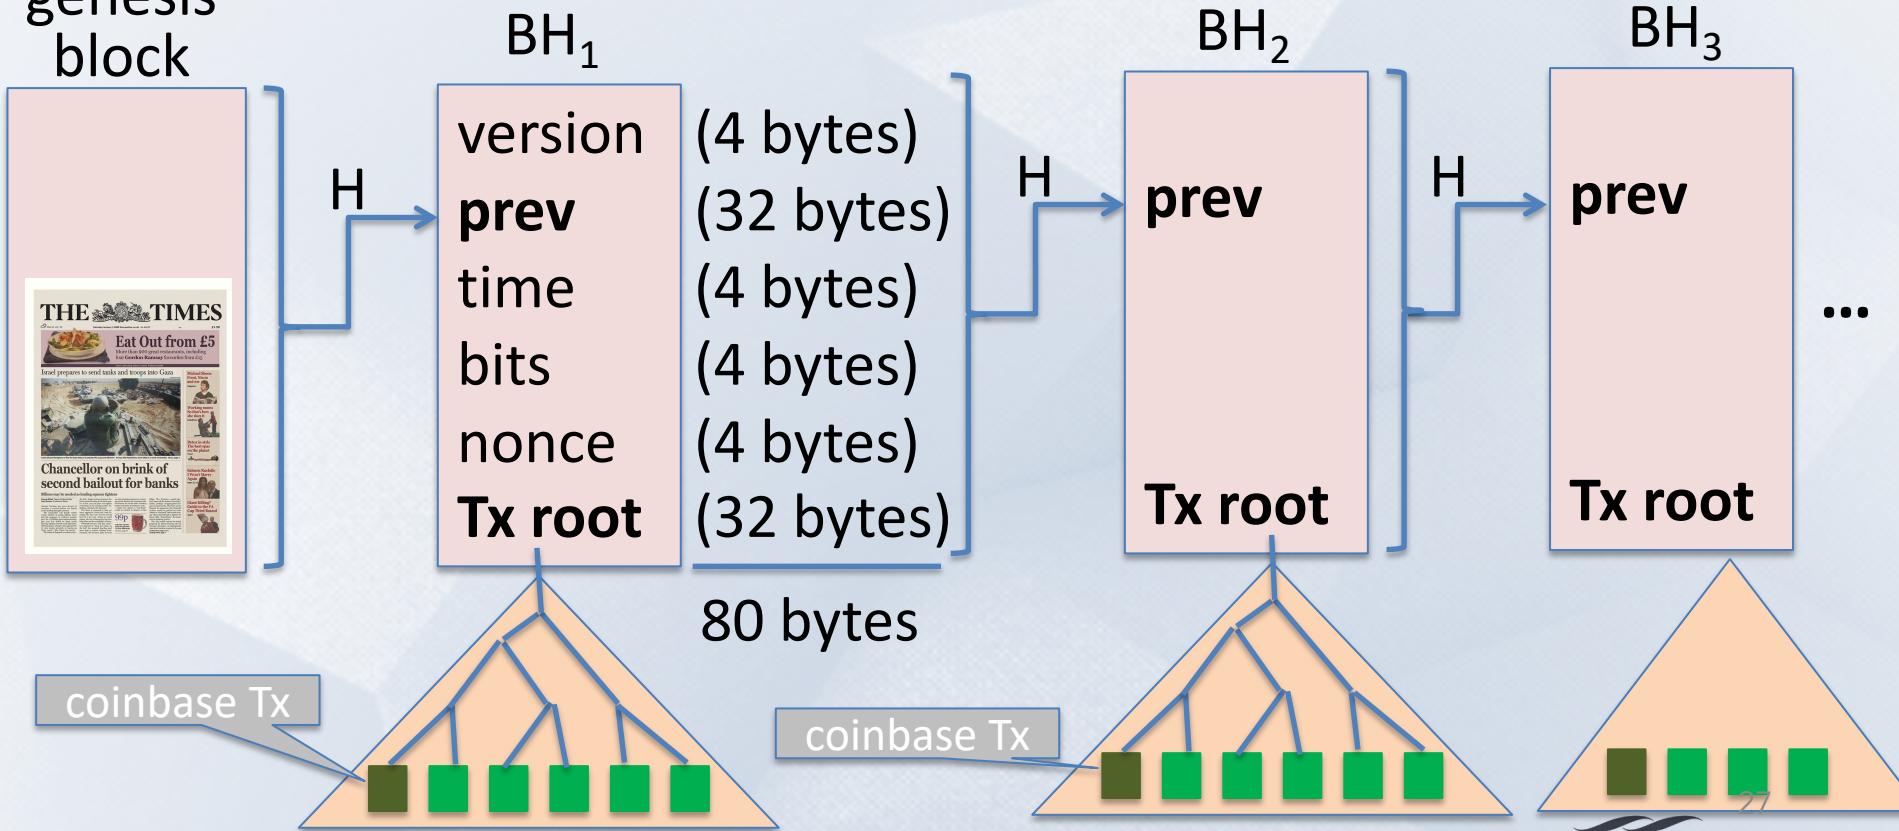
Liveness:

- to block a Tx from being posted, need to convince 51% of mining power **

(some sub 50% censorship attacks, such as feather forks)

Bitcoin blockchain: a sequence of block headers, 80 bytes each

genesis
block



Bitcoin blockchain: a sequence of block headers, 80 bytes each

time: time miner assembled the block. Self reported.
(block rejected if too far in past or future)

bits: proof of work difficulty

nonce: proof of work solution



} for choosing a leader (next week)

Merkle tree: payer can give a short proof that Tx is in the block

new block every \approx 10 minutes.

An example

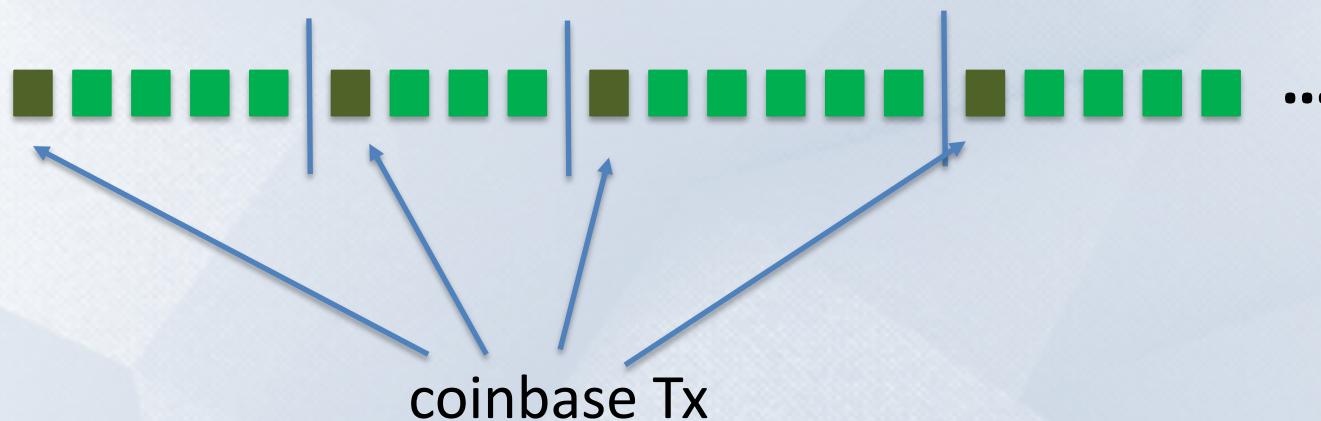
Height	Mined	Miner	Size	Tx data	#Tx
648494	17 minutes	Unknown	1,308,663 bytes		1855
648493	20 minutes	SlushPool	1,317,436 bytes		2826
648492	59 minutes	Unknown	1,186,609 bytes		1128
648491	1 hour	Unknown	1,310,554 bytes		2774
648490	1 hour	Unknown	1,145,491 bytes		2075
648489	1 hour	Poolin	1,359,224 bytes		2622

Block 648493

Timestamp	2020-09-15 17:25	
Height	648493	
Miner	SlushPool	(from coinbase Tx)
Number of Transactions	2,826	
Difficulty (D)	17,345,997,805,929.09	(adjusts every two weeks)
Merkle root	350ccb917c918774c93e945b960a2b3ac1c8d448c2e67839223bbcf595baff89	
Transaction Volume	11256.14250596 BTC	
Block Reward	6.25000000 BTC	
Fee Reward	0.89047154 BTC	(Tx fees given to miner in coinbase Tx)

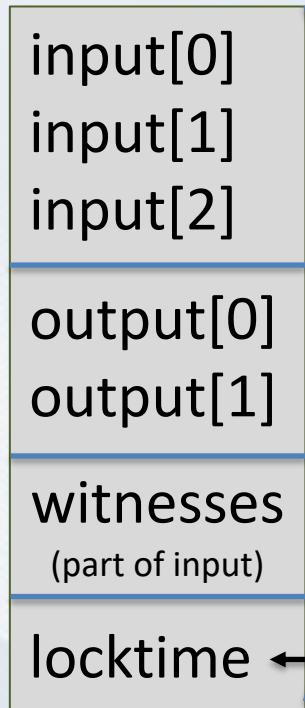
This lecture

View the blockchain as a sequence of Tx (append-only)

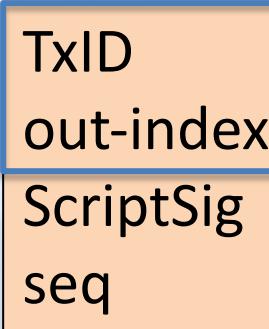


Tx structure (non-coinbase)

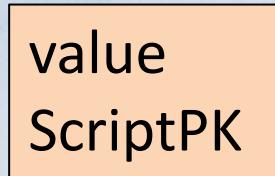
inputs
outputs
(segwit)
(4 bytes)



input:



output:

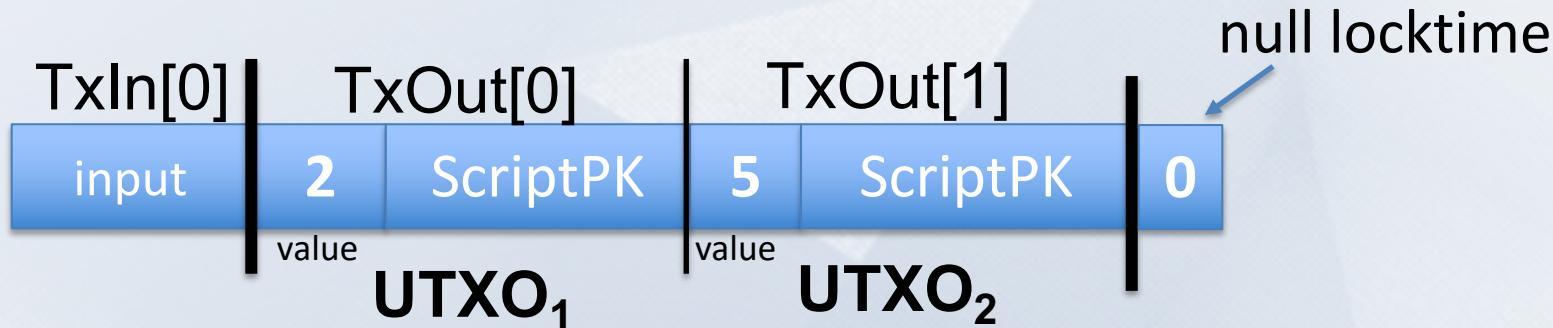


$$\#BTC = \text{value}/10^8$$

earliest block # that can include Tx

Example

Tx1:
(funding Tx)



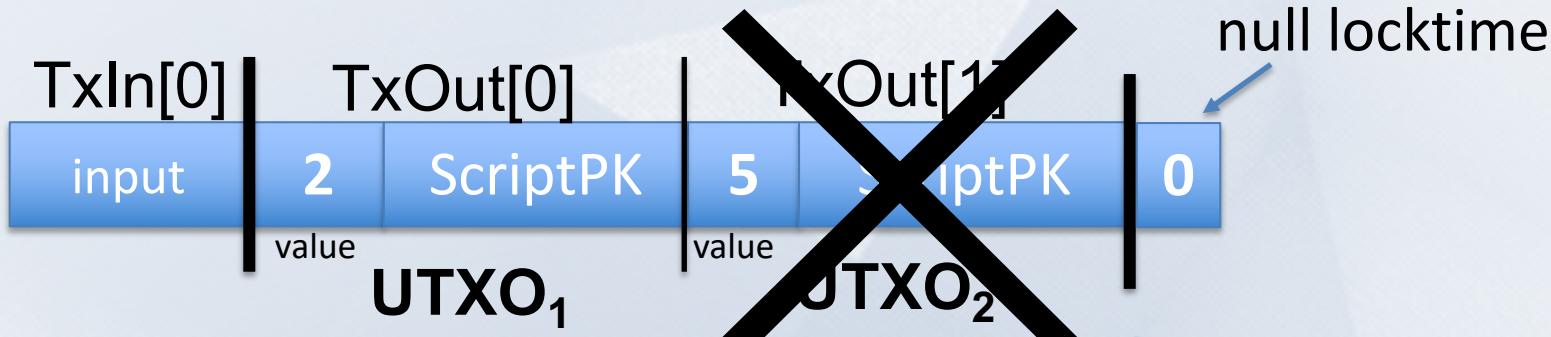
UTXO: unspent Tx output

Tx2:
(spending Tx)



Example

Tx1:
(funding Tx)



UTXO: unspent Tx output

Tx2:
(spending Tx)



Validating Tx2

Miners check (for each input):

1. The program ScriptSig | ScriptPK returns true
2. TxID | index is in the current UTXO set
3. sum input values \geq sum output values

After Tx2 is posted, miners remove UTXO₂ from UTXO set

program from funding Tx:
under what conditions
can UTXO be spent



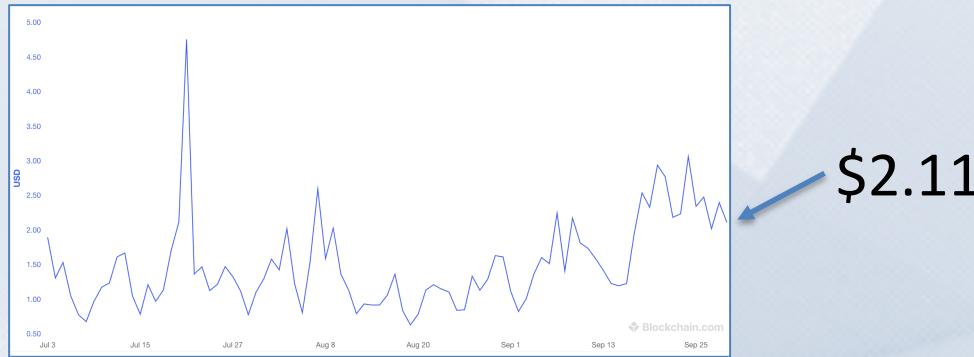
An example (block 648493) [2826 Tx]

COINBASE (Newly Generated Coins)	→	1CK6KHY6MHgYvmRQ4PAafKYDrg1ejbH1cE OP_RETURN OP_RETURN	7.14047154 BTC	🌐
Tx0			0.00000000 BTC	
			6.25 + Tx fees =	7.14047154 BTC
3PuJbxJS1pKxf8EdVR18yBkD1fPAbgUtyw input	0.72333974 BTC (input UTXO value)	→	1E5Ao1VUnA5BhffvXf2Xmud6avUgwkFnJv bc1qr8k3e0vx06lpu3j7m858pa2ak9tyr56ttwvefk bc1qdrxve8kua3yz5dgx6wf3u95ngh0d3e648... 14ZhjuXpQ5jCDjtAy7ZMu3hfEQCWewzLw7	0.00917379 BTC 0.61504199 BTC 0.09290152 BTC 0.00616444 BTC
Tx1	0.00005800 BTC (Tx fee)		outputs	0.72328174 BTC
17MWze4Z1uP1jnvqvj7SAnGtxcoVq11H8A Tx2	0.05000000 BTC 0.00192000 BTC (Tx fee)	→	3G3C2RFQ8gsf77EQpdR4ZReChWFKEHhxVU	0.04808000 BTC
				0.04808000 BTC

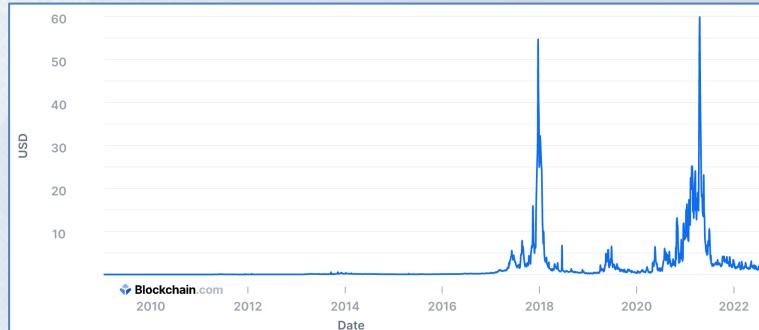
sum of fees in block added to coinbase Tx

Tx fees

Bitcoin average Tx fees in USD (last 60 days, sep. 2023)



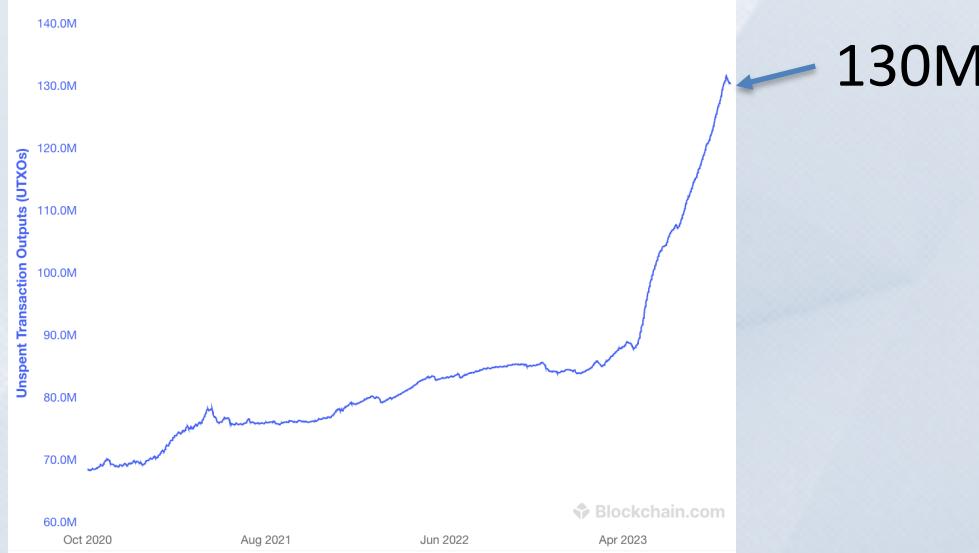
Bitcoin average Tx fees in USD (all time)



All value in Bitcoin is held in UTXOs

Unspent Transaction Outputs

The total number of valid unspent transaction outputs. This excludes invalid UTXOs with opcode OP_RETURN



Sep. 2023: miners need to store $\approx 130M$ UTXOs in memory

Focusing on Tx2: TxIn[0]

Value	0.05000000 BTC
Pkscript	OP_DUP OP_HASH160 45b21c8a0cb687d563342b6c729d31dab58e3a4e OP_EQUALVERIFY OP_CHECKSIG
Sigscript	304402205846cace0d73de82dfbdeba4d65b9856d7c1b1730eb401cf4906b2401a69b dc90220589d36d36be64e774c8796b96c011f29768191abeb7f56ba20ffb0351280860 c01 03557c228b080703d52d72ead1bd93fc72f45c4ddb4c2b7a20c458e2d069c8dd9e

from UTXO
(Bitcoin script)

from TxIn[0]

Bitcoin Script

A stack machine. Not Turing Complete: no loops.

Quick survey of op codes:

1. OP_TRUE (OP_1), OP_2, ..., OP_16: push value onto stack

81

82

96

2. OP_DUP: push top of stack onto stack

118

Bitcoin Script

3. control:

99 OP_IF <statements> OP_ELSE <statements> OP_ENDIF

105 OP_VERIFY: abort fail if top = false

106 OP_RETURN: abort and fail

what is this for? ScriptPK = [OP_RETURN, <data>]

136 OP_EQVERIFY: pop, pop, abort fail if not equal

Bitcoin Script

4. arithmetic:

OP_ADD, OP_SUB, OP_AND, ...: pop two items, add, push

5. crypto:

OP_SHA256: pop, hash, push

OP_CHECKSIG: pop pk, pop sig, verify sig. on Tx, push 0 or 1

6. Time: OP_CheckLockTimeVerify (CLTV):

fail if value at the top of stack > Tx locktime value.

usage: UTXO can specify min-time when it can be spent

Example: a common script

```
<sig> <pk> DUP HASH256 <pkhash> EQVERIFY CHECKSIG
```

stack: empty

init

<sig> <pk>

push values

<sig> <pk> <pk>

DUP

<sig> <pk> <hash>

HASH256

<sig> <pk> <hash> <pkhash>

push value

<sig> <pk>

EQVERIFY

1

CHECKSIG

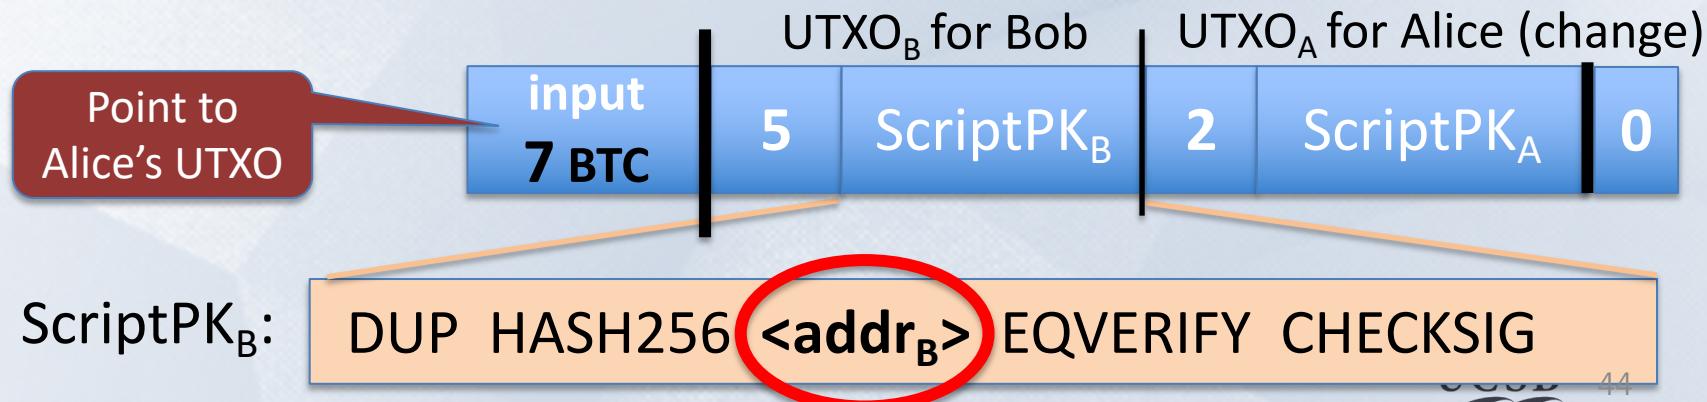
⇒ successful termination

verify(pk, Tx, sig)

Transaction types: (1) P2PKH

Alice want to pay Bob 5 BTC: pay to public key hash

- step 1: Bob generates sig key pair $(pk_B, sk_B) \leftarrow Gen()$
- step 2: Bob computes his Bitcoin address as $addr_B \leftarrow H(pk_B)$
- step 3: Bob sends $addr_B$ to Alice
- step 4: Alice posts Tx:

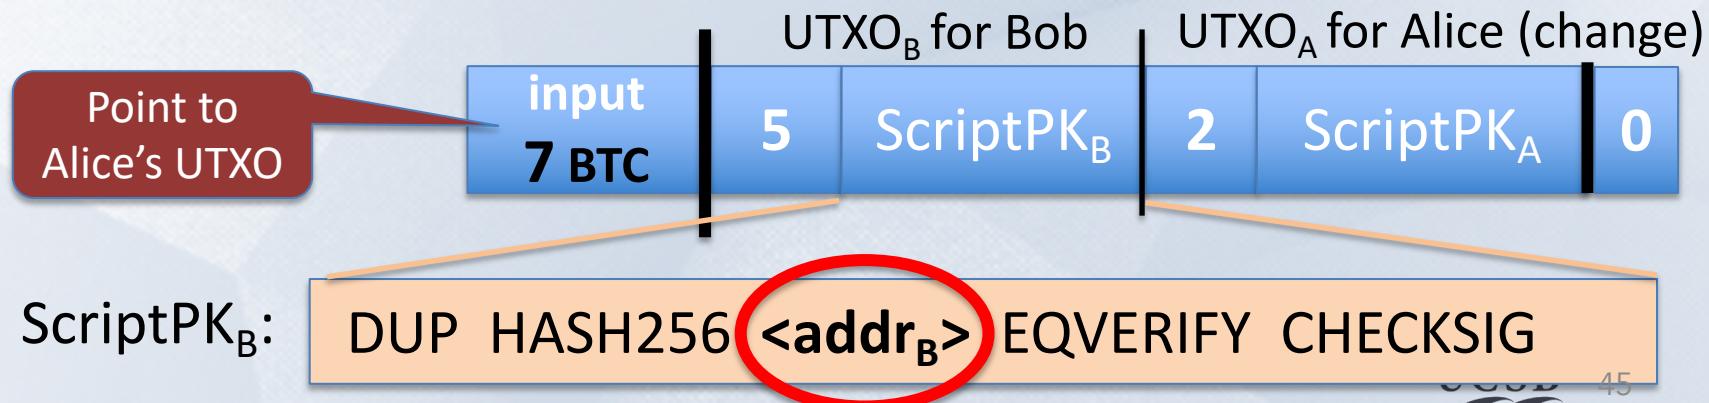


Transaction types: (1) P2PKH

pay to public key hash

“input” contains ScriptSig that authorizes spending Alice’s UTXO

- example: ScriptSig contains Alice’s signature on Tx
⇒ miners cannot change ScriptPK_B (will invalidate Alice’s signature)



Transaction types: (1) P2PKH

Later, when Bob wants to spend his UTXO: create a Tx_{spend}

Tx_{spend} :



points to
UTXO_B

$<\text{sig}>$ $<\text{pk}_B>$

(authorizes spending UTXO_B)

$<\text{sig}> = \text{Sign}(\text{sk}_B, \text{Tx})$ where $\text{Tx} = (Tx_{spend} \text{ excluding all ScriptSigs})$ (SIGHASH_ALL)

Miners validate that $\text{ScriptSig}_B \mid \text{ScriptPK}_B$ returns true

P2PKH: comments

- Alice specifies recipient's pk in UTXO_B
- Recipient's pk is not revealed until UTXO is spent
(some security against attacks on pk)
- Miner cannot change $\langle \text{Addr}_B \rangle$ and steal funds:
invalidates Alice's signature that created UTXO_B

Segregated Witness

ECDSA malleability:

- Given (m, sig) anyone can create (m, sig') with $\text{sig} \neq \text{sig}'$
- ⇒ miner can change sig in Tx and change TxID = SHA256(Tx)
 - ⇒ Tx issuer cannot tell what TxID is, until Tx is posted
 - ⇒ leads to problems and attacks

Segregated witness: signature is moved to witness field in Tx
TxID = Hash(Tx without witnesses)

Transaction types: (2) P2SH: pay to script hash

(pre SegWit in 2017)

Let's payer specify a redeem script (instead of just pkhash)

Usage: payee publishes hash(redeem script) ← Bitcoin
addr.

payer sends funds to that address

ScriptPK in UTXO: HASH160 <H(redeem script)> EQUAL

ScriptSig to spend: <sig₁> <sig₂> ... <sig_n> <redeem script>

payer can specify complex conditions for when UTXO can be spent

P2SH

Miner verifies:

- (1) <ScriptSig> $\text{ScriptPK} = \text{true}$ \leftarrow payee gave correct script
- (2) $\text{ScriptSig} = \text{true}$ \leftarrow script is satisfied

Example P2SH: multisig

Goal: spending a UTXO requires t-out-of-n signatures

Redeem script for 2-out-of-3: (set by payer)

<2> <PK₁> <PK₂> <PK₃> <3> CHECKMULTISIG

 hash gives P2SH address

ScriptSig to spend: (by payee)

<0> <sig1> <sig3> <redeem script>

END OF LECTURE

Next lecture: interesting scripts,
wallets, and how to manage crypto assets

