

Consensus in the Internet Setting

Yu Feng

University of California, Santa Barbara



Slides adapted from Stanford CS251

Recap of the Last Lecture

- Byzantine Generals Problem
- Definition of Byzantine adversary
 - **Byzantine:** Adversarial nodes can deviate from the protocol arbitrarily!
- Synchronous and asynchronous networks
 - **Synchronous network:** known upper bound Δ on network delay
- Byzantine Broadcast
- Dolev-Strong (1983)
- State Machine Replication (SMR)
- Security properties for SMR protocols: Safety and Liveness

Sybil Attack

How to select the nodes that participate in consensus?



Two variants:

- *Permissioned*: There is a *fixed* set of nodes (previous lecture).
- *Permissionless*: Anyone is free to join the protocol at any time.

Can we accept any node that has a signing key to participate in consensus?

Sybil Attack!

Sybil Attack

How to select the nodes that participate in consensus?



Two variants:

- *Permissioned*: There is a *fixed* set of nodes (previous lecture).
- *Permissionless*: Anyone is free to join the protocol at any time.

Can we accept any node that has a signing key to participate in consensus?

In a **sybil attack**, a single adversary impersonates many different nodes, outnumbering the honest nodes and potentially disrupting consensus.

Sybil Resistance

Consensus protocols with Sybil resistance are typically based on a bounded (scarce) resource:

	Resource dedicated to the protocol	Some Example Blockchains
Proof-of-Work	Total computational power	Bitcoin, PoW Ethereum...
Proof-of-Stake	Total number of coins	Algorand, Cardano, Cosmos, PoS Ethereum...
Proof-of-Space/Time	Total storage across time	Chia, Filecoin...

We assume that the adversary controls a small fraction of the scarce resource!

Resource gives the power to influence the protocol.

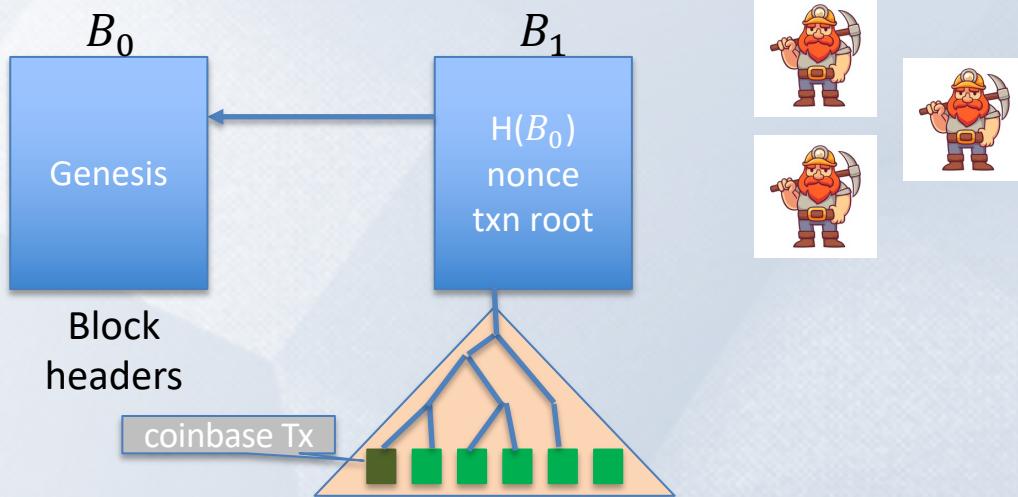
Adversary has less influence than honest nodes.

Bitcoin: Mining

To mine a new block, a miner must find *nonce* such that

$$H(h_{prev}, \text{txn root}, \text{nonce}) < \text{Target} = \frac{2^{256}}{D}$$

Each miner tries different nonces until one of them finds a nonce that satisfies the above equation.



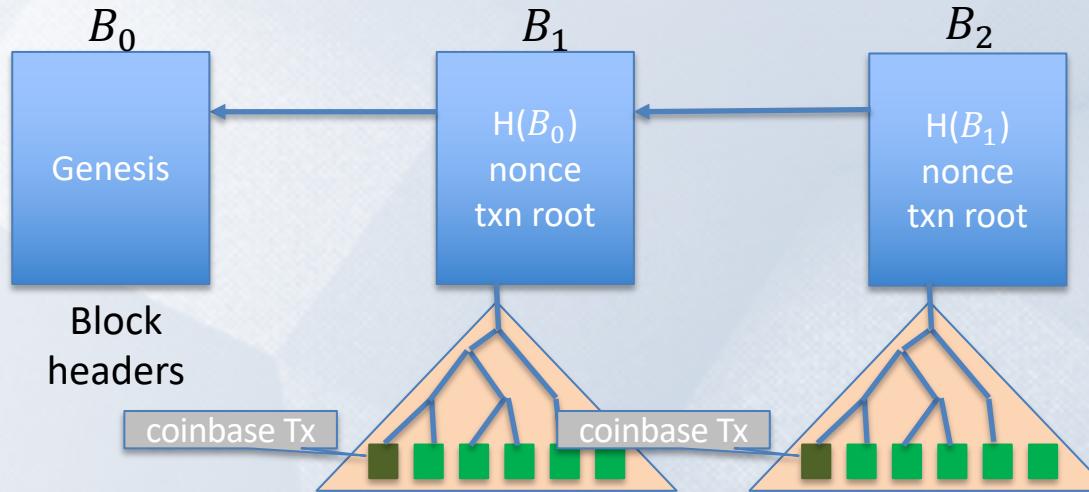
Bitcoin: Mining

To mine a new block, a miner must find *nonce* such that

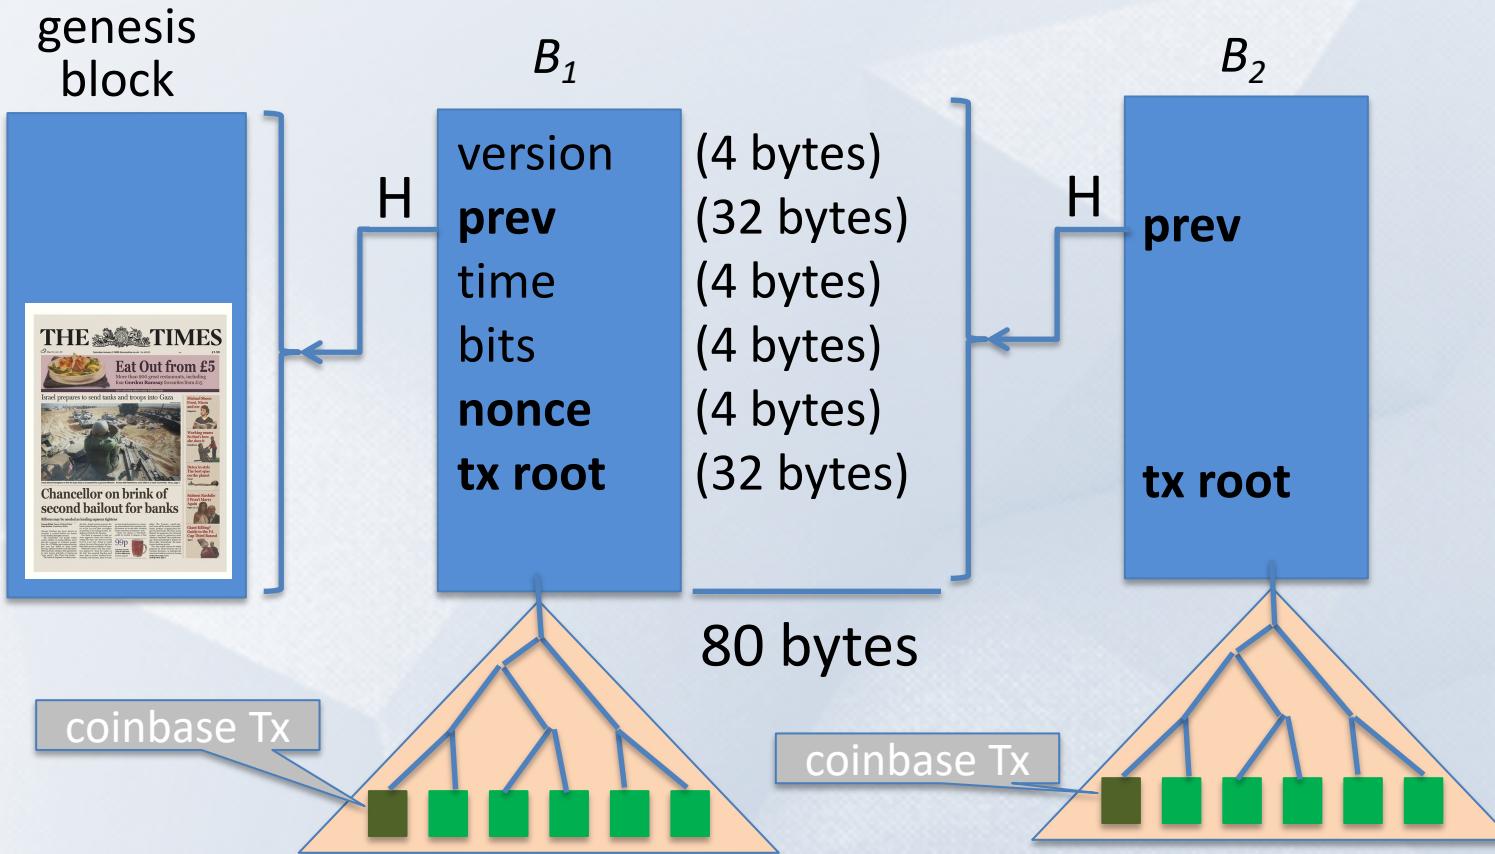
$$H(h_{prev}, \text{txn root}, \text{nonce}) < \text{Target} = \frac{2^{256}}{D}$$

Difficulty: How many nonces on average miners try until finding a block?

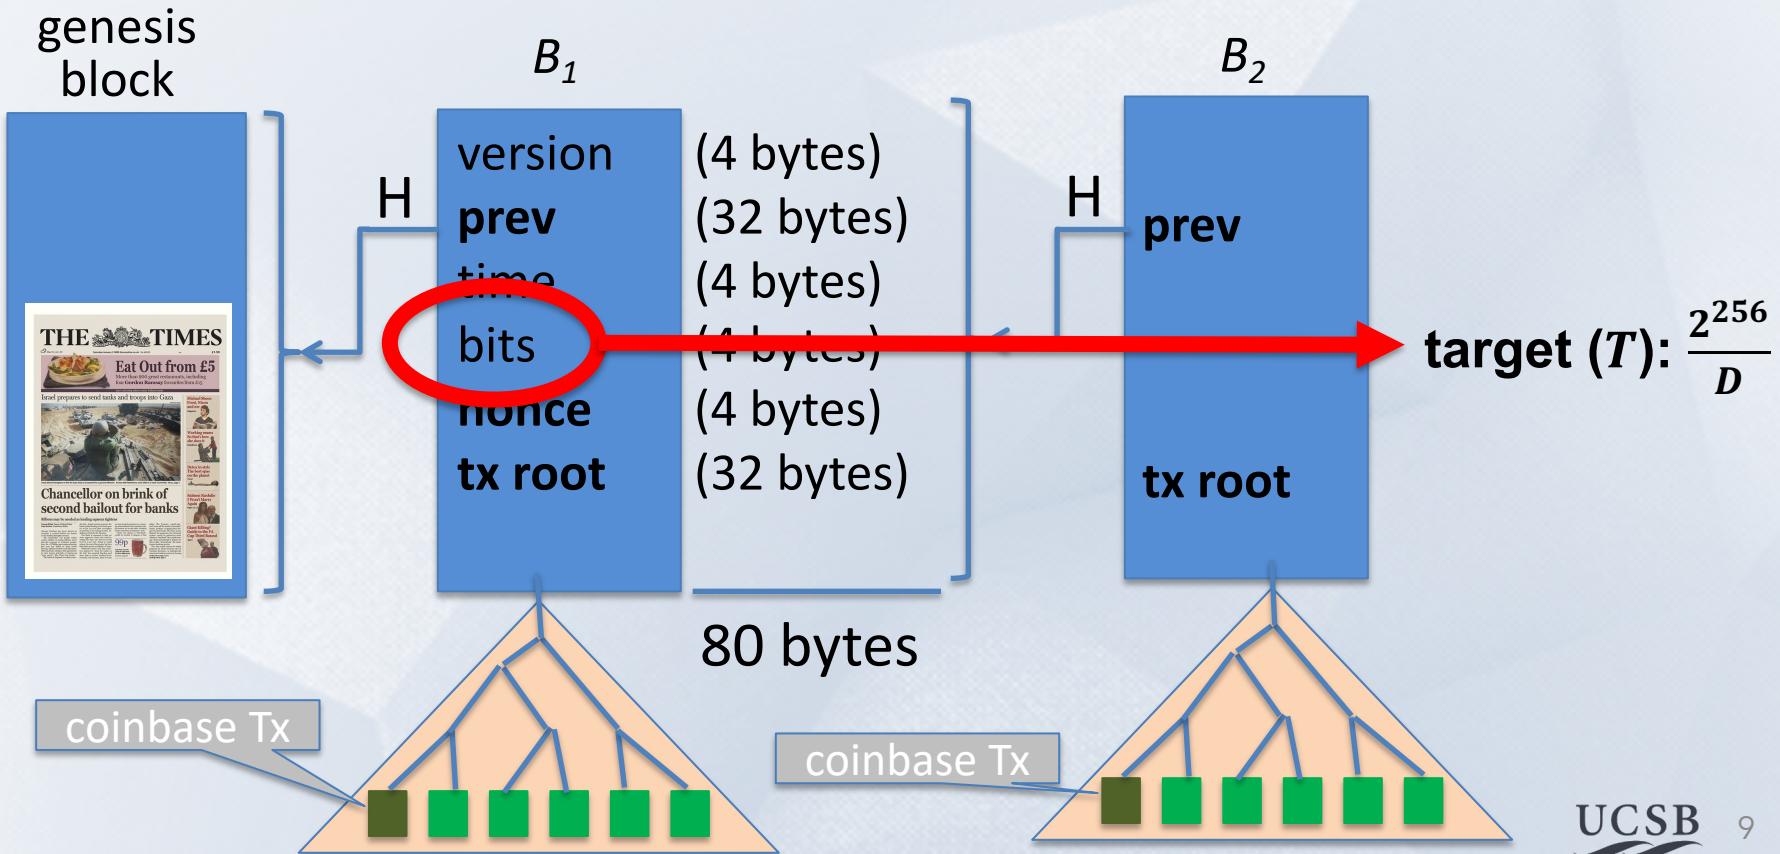
Each miner tries different nonces until one of them finds a nonce that satisfies the above equation.



Bitcoin: Mining



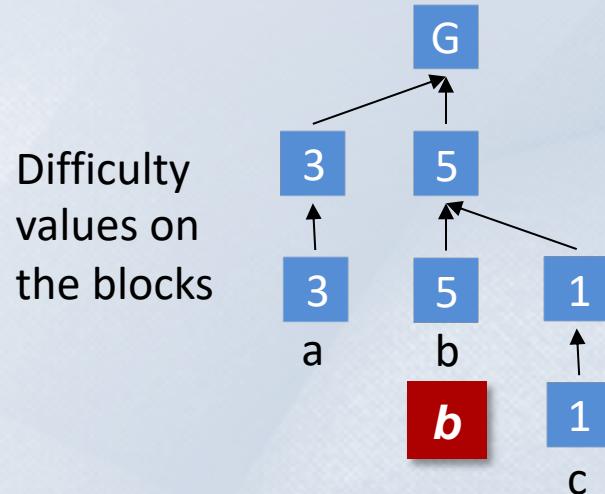
Bitcoin: Mining



Nakamoto Consensus

Bitcoin uses Nakamoto consensus:

- **Fork-choice / proposal rule:** At any given time, each honest miner attempts to extend (i.e., mines on the tip of) the heaviest chain *held* in its view (Ties broken adversarially).

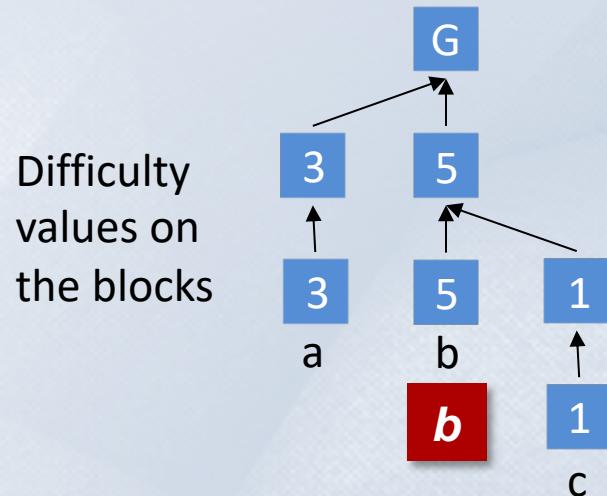


Nakamoto Consensus

Chain with the highest difficulty, i.e, largest sum of the difficulty D within blocks!

Bitcoin uses Nakamoto consensus:

- Fork-choice / proposal rule: At any given time, each honest miner attempts to extend (i.e., mines **on the tip of**) the heaviest chain *held* in its view (Ties broken adversarially).

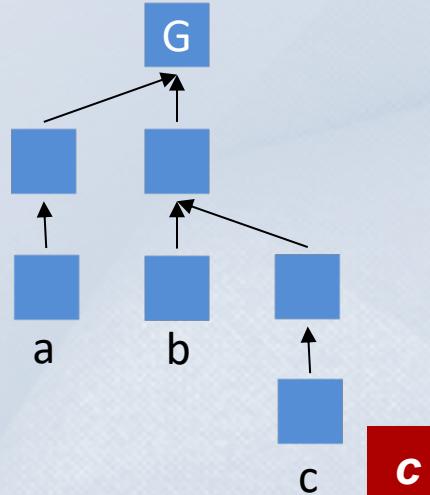


Nakamoto Consensus

Chain with the highest difficulty, i.e, largest sum of the difficulty D within blocks!

Bitcoin uses Nakamoto consensus:

- Fork-choice / proposal rule: At any given time, each honest miner attempts to extend (i.e., mines **on the tip of**) the heaviest (longest for us) chain held in its view (Ties broken adversarially).



Nakamoto Consensus

Bitcoin uses Nakamoto consensus:

- **Fork-choice / proposal rule:** At any given time, each honest miner attempts to extend (i.e., mines on the tip of) the heaviest (longest for us) chain *held* in its view (Ties broken adversarially).
- **Confirmation rule:** Each miner confirms the block (along with its prefix) that is k -deep within the longest chain in its view.
 - In practice, $k = 6$.
 - Miners and clients accept the transactions in the latest confirmed block and its prefix as their log.
 - Note that *confirmation* is **different** from *finalization*.
- **Leader selection rule:** Proof-of-Work.

Nakamoto Consensus

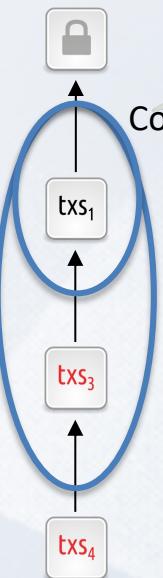
Chain with the highest difficulty, i.e, largest sum of the difficulty D within blocks!

Bitcoin uses Nakamoto consensus:

- **Fork-choice / proposal rule:** At any given time, each honest miner attempts to extend (i.e., mines on the tip of) the heaviest (longest for us) chain *held* in its view (Ties broken adversarially).
- **Confirmation rule:** Each miner confirms the block (along with its prefix) that is k -deep within the longest chain in its view.
 - In practice, $k = 6$.
 - Miners and clients accept the transactions in the latest confirmed block and its prefix as their log.
 - Note that confirmation is **different** from finalization.
- **Leader selection rule:** Proof-of-Work.

Nakamoto Consensus

$k=2$



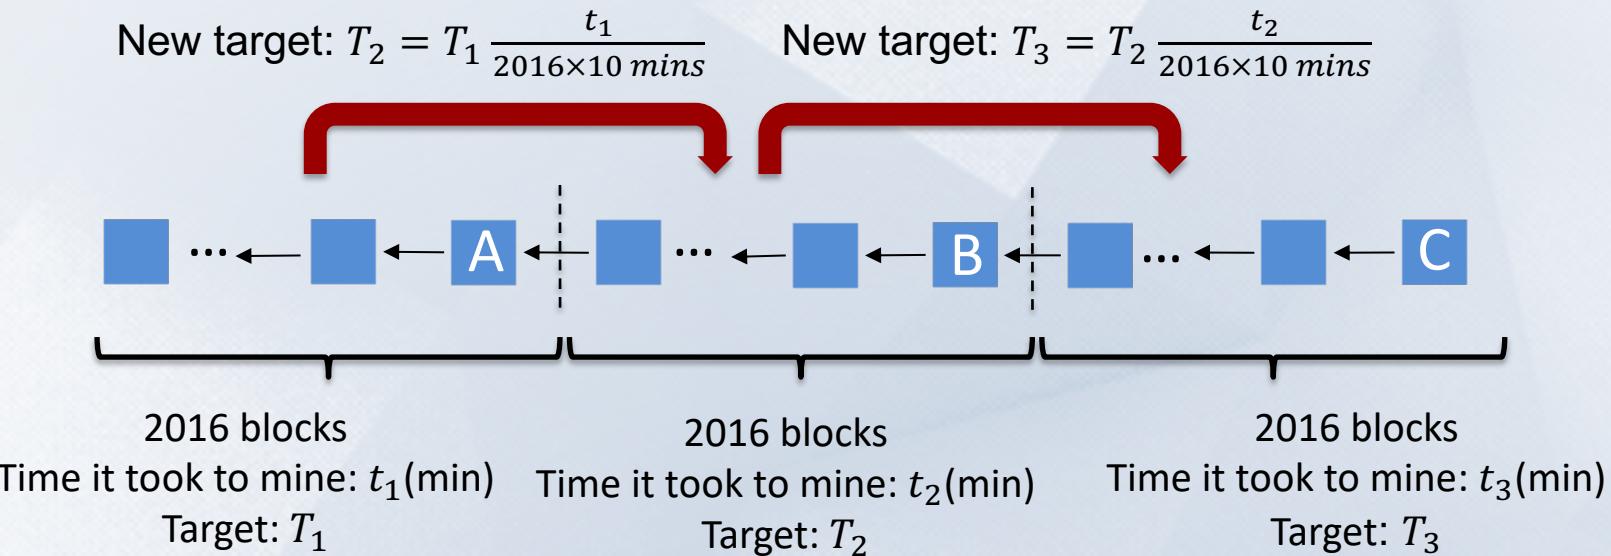
Confirmed



Alice's log: $txs_1 \ txs_3$



Bitcoin: Difficulty Adjustment



New target is not allowed to be more than 4x old target.
New target is not allowed to be less than $\frac{1}{4}$ x old target.

t_2 : difference between the timestamps in B and A
 t_3 : difference between the timestamps in C and B

Consensus in the Internet Setting

Characterized by *open participation*.

Challenges:

- Adversary can create many Sybil nodes to take over the protocol.
- Honest nodes can come and go at will.

Requirements:

Achieved by Bitcoin!

- Limit adversary's participation.
 - **Sybil resistance (e.g., Proof-of-Work)!**
- Maintain availability (liveness) of the protocol when the honest nodes come and go at will, resulting in changes in the number of nodes.
 - **Dynamic availability!**

Security?

Can we show that Bitcoin is a secure state machine replication (SMR) protocol (satisfies safety and liveness) under synchrony against a Byzantine adversary?

$$\beta(t)$$

$\in [0,1]$ for all t

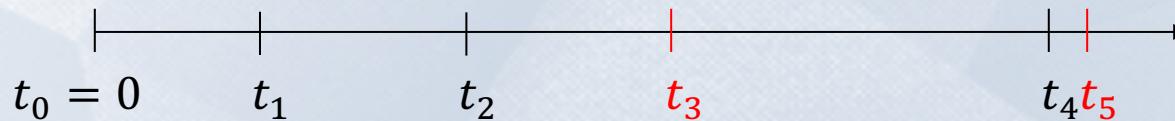


**Fraction of the mining power
controlled by the adversary at time t .**

What is the highest $\beta(t)$ for which Bitcoin is secure??

Model for Bitcoin

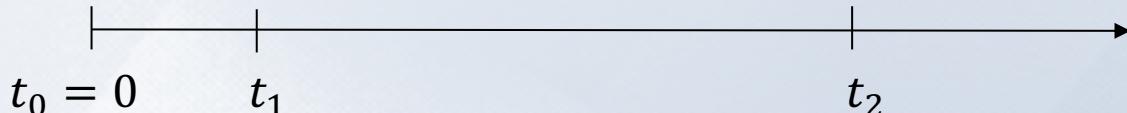
- Many different miners, each with *infinitesimal* power.
Total mining rate (growth rate of the chain): λ (1/minutes). In Bitcoin, $\lambda = 1/10$.
- Suppose **Adversary** is Byzantine and controls $\beta < \frac{1}{2}$ fraction of the mining power.
 - Adversarial mining rate: $\lambda_a = \beta\lambda$
 - Honest mining rate: $\lambda_h = (1 - \beta)\lambda$
- Network is **synchronous** with a known upper bound Δ on delay.



Reminder: Why is safety important?

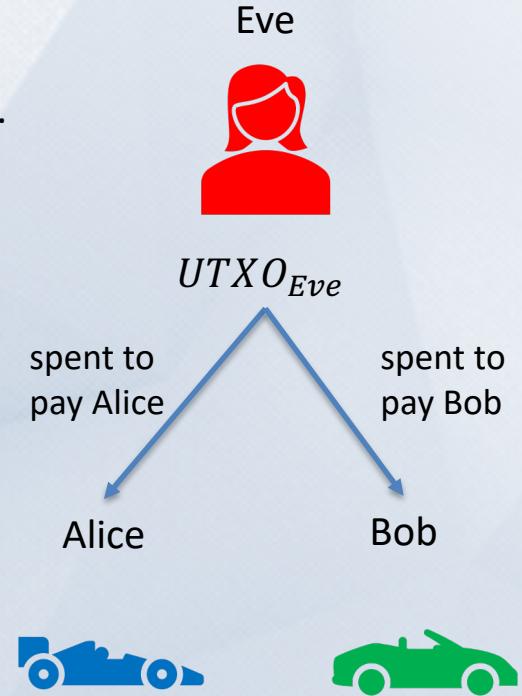
Suppose Eve has a UTXO.

- tx_1 : transaction spending Eve's UTXO to pay to car vendor Alice.
- tx_2 : transaction spending Eve's UTXO to pay to car vendor Bob.



- Alice's ledger at time t_1 contains tx_1 :
 $LOG_{t_1}^{Alice} = \langle tx_1 \rangle$
- Alice thinks it received Eve's payment and sends over the car.

- Bob's ledger at time t_2 contains tx_2 :
 $LOG_{t_2}^{Bob} = \langle tx_2 \rangle$
- Bob thinks it received Eve's payment and sends over the car.



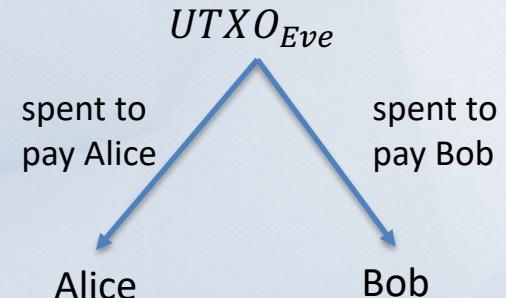
Reminder: Why is safety important?

Suppose Eve has a UTXO.

- tx_1 : transaction spending Eve's UTXO to pay to car vendor Alice.
- tx_2 : transaction spending Eve's UTXO to pay to car vendor Bob.

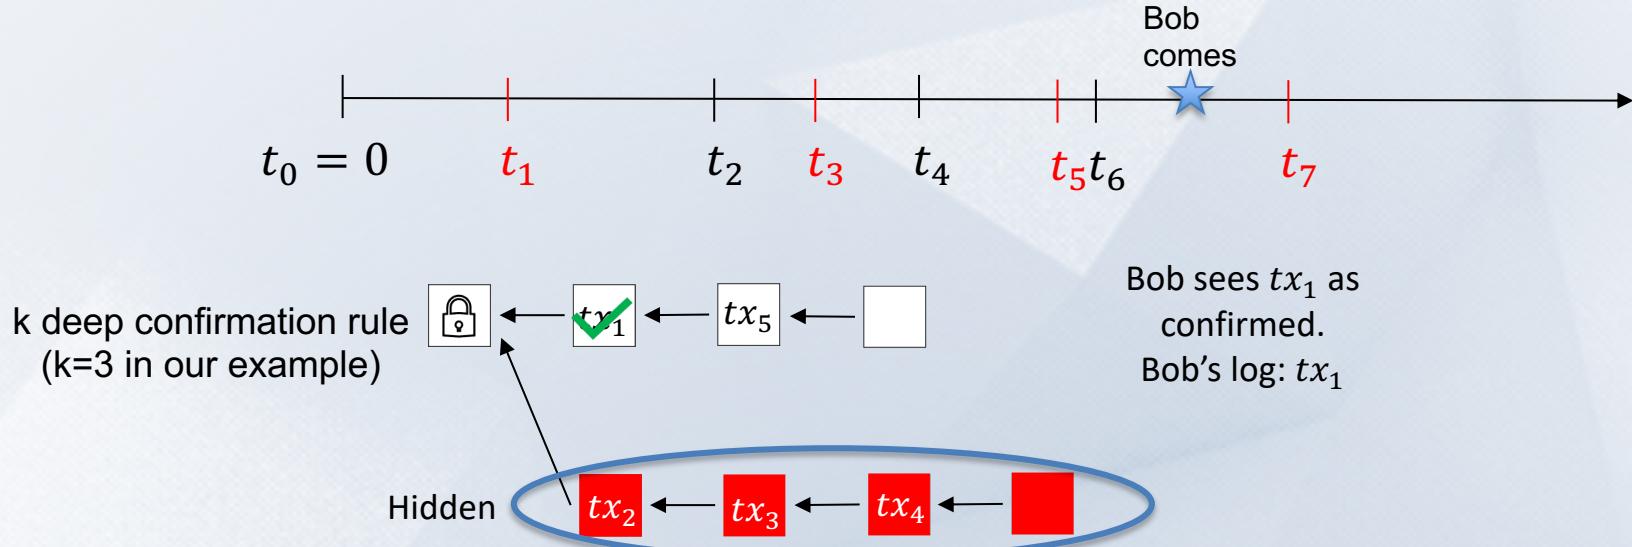


- Alice's ledger at time t_1 contains tx_1 :
 $LOG_{t_1}^{Alice} = \langle tx_1 \rangle$
- Alice thinks it received Eve's payment and sends over the car.
- Bob's ledger at time t_2 contains tx_2 :
 $LOG_{t_2}^{Bob} = \langle tx_2 \rangle$
- Bob thinks it received Eve's payment and sends over the car.



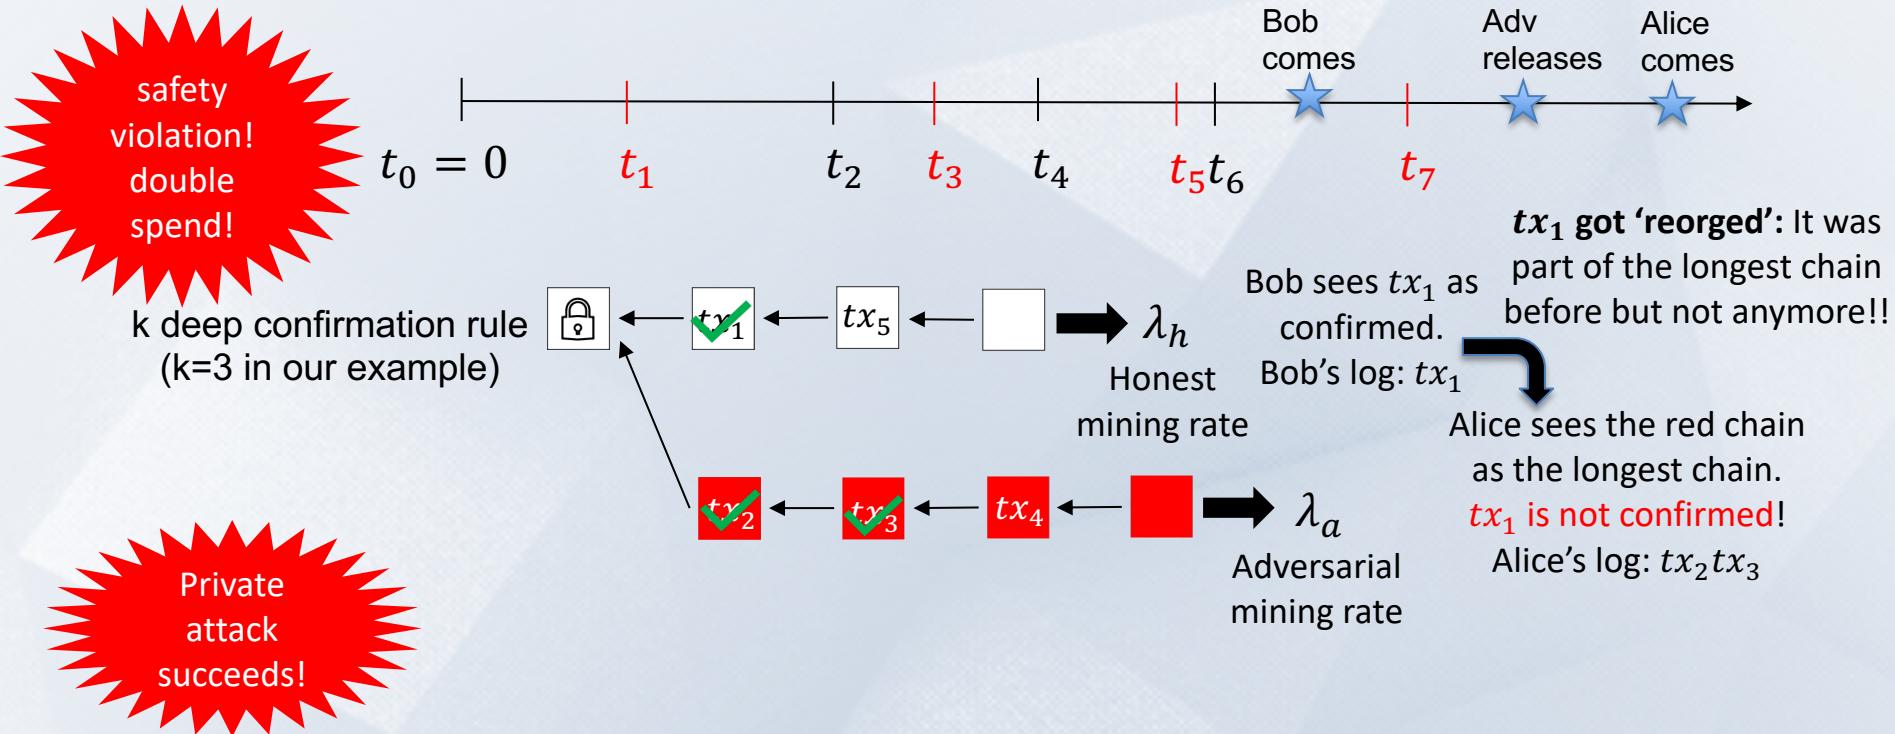
When safety is violated, Eve can double-spend!

Nakamoto's Private Attack: $\beta \geq 1/2$

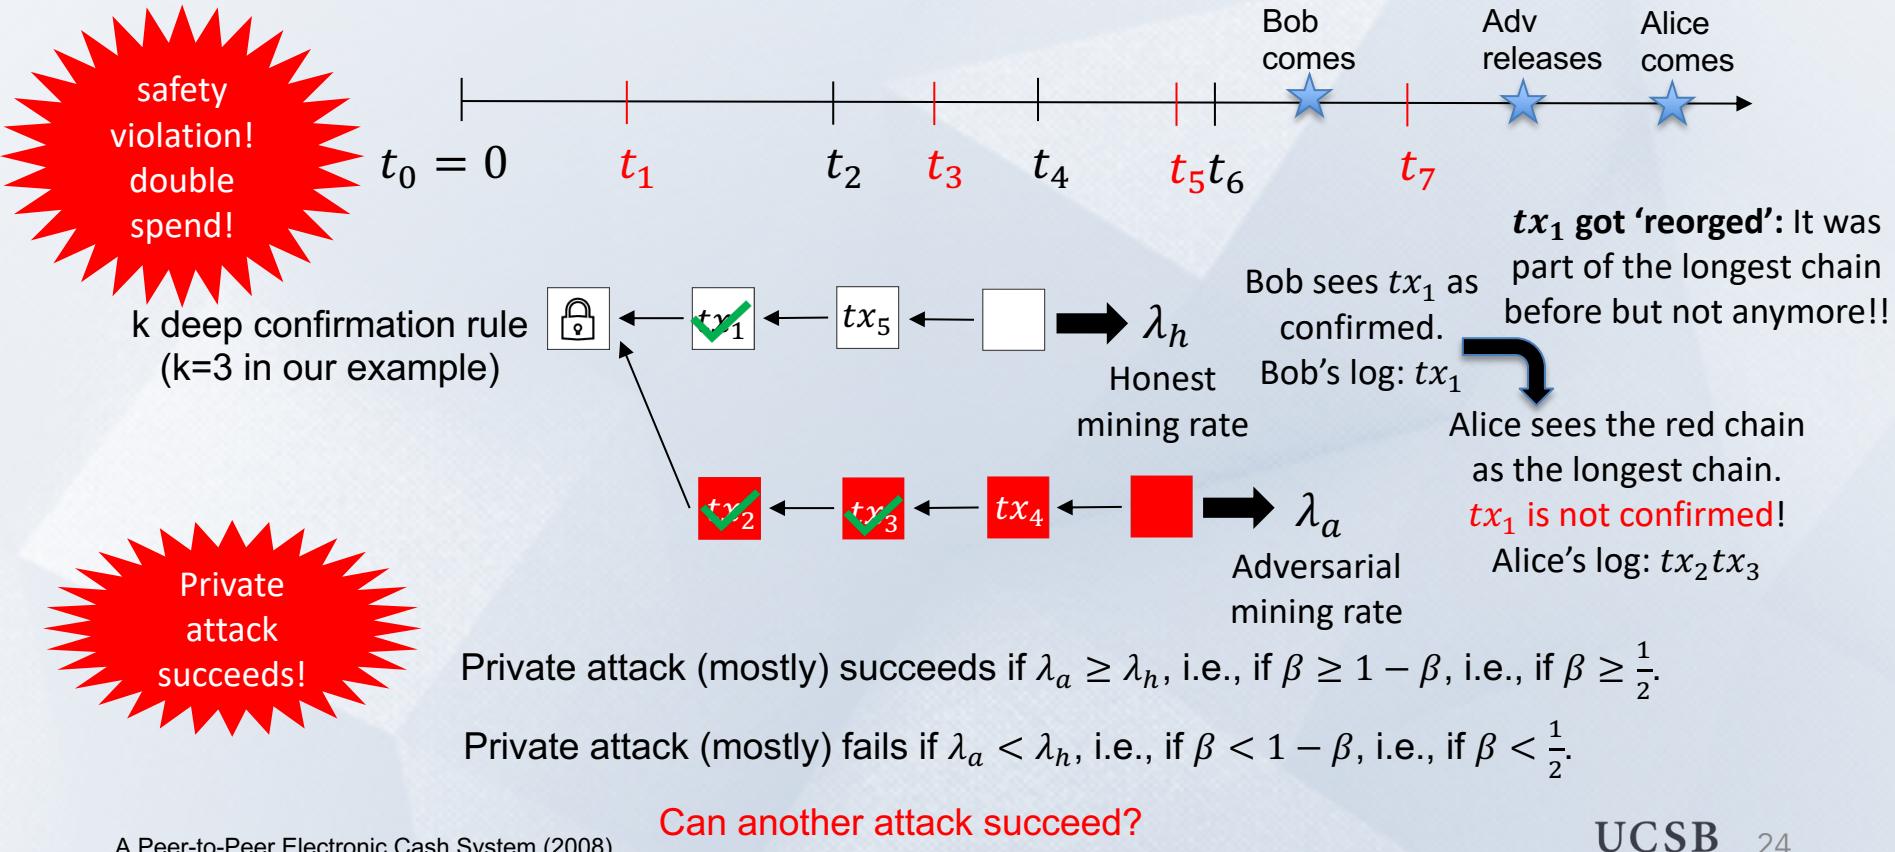


Let's show that Bitcoin is insecure if $\beta(t) \geq 1/2$

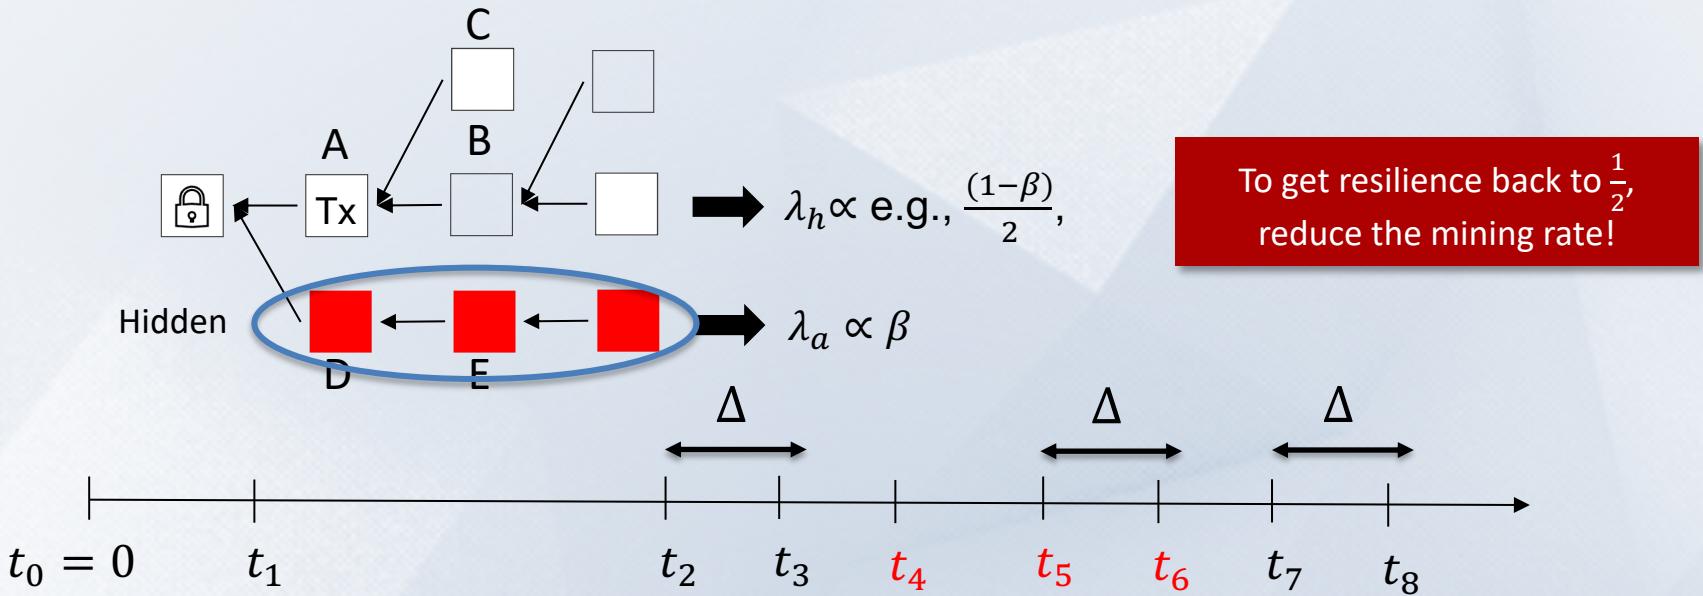
Nakamoto's Private Attack: $\beta \geq 1/2$



Nakamoto's Private Attack: $\beta \geq 1/2$



Forking



Multiple honest blocks at the same height due to network delay.
Adversary's chain grows at rate proportional to (shown by \propto) β !
Honest miners' chain grows at rate less than $1 - \beta$ because of forking!

Now, adversary succeeds if $\beta \geq \frac{(1-\beta)}{2}$, which implies $\beta \geq \frac{1}{3}!!$

Reminder for SMR Security

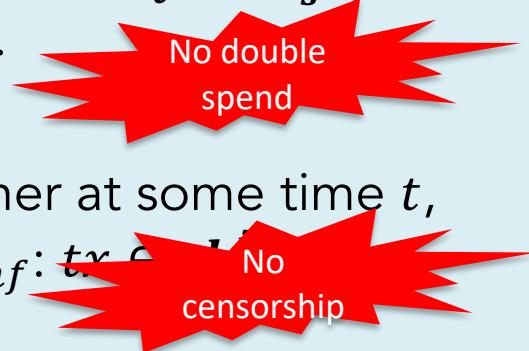
Let's recall the security definition for state machine replication (SMR) protocols. Let ch_t^i denote the confirmed (i.e., k -deep) of a client i at time t .

Safety (Consistency):

- For any two clients i and j , and times t and s : $ch_t^i \leq ch_s^j$ (prefix of) or vice versa, i.e., chains are consistent.

Liveness:

- If a transaction tx is input to an honest miner at some time t , then for all clients i , and times $s \geq t + T_{conf}$: $tx \in ch_s^i$



Security Theorem

Theorem: If $\beta < 1/2$, there exists a small enough mining rate $\lambda(\Delta, \beta) = \lambda_a + \lambda_h$ such that Bitcoin satisfies safety and liveness except with error probability $\epsilon \equiv e^{-\Omega(k)}$ under synchronous network (recall that k is used in the k deep confirmation rule).

- $e^{-\Omega(k)}$ is the error probability for confirmation.
- Latest result for bounding the error probability as a function of k :

$$\epsilon \leq \left(2 + 2 \sqrt{\frac{1-\beta}{\beta}} \right) (4\beta(1-\beta))^k$$

- We say ‘confirmation’ instead of finalization because when you *confirm* a block or transaction, you *confirm* it with an error probability...
- ...unlike *finalizing* a block where there is no error probability*.

The Bitcoin Backbone Protocol: Analysis and Applications (2015)
Analysis of the Blockchain Protocol in Asynchronous Networks (2016)
Analysis of Nakamoto Consensus (2019)
Everything is a Race and Nakamoto Always Wins (2022)
Bitcoin’s Latency–Security Analysis Made Simple (2022)

Security Theorem

Theorem: If $\beta < 1/2$, there exists a small enough mining rate $\lambda(\Delta, \beta) = \lambda_a + \lambda_h$ such that Bitcoin satisfies safety and liveness except with error probability $\epsilon \equiv e^{-\Omega(k)}$ under synchronous network (recall that k is used in the k deep confirmation rule).

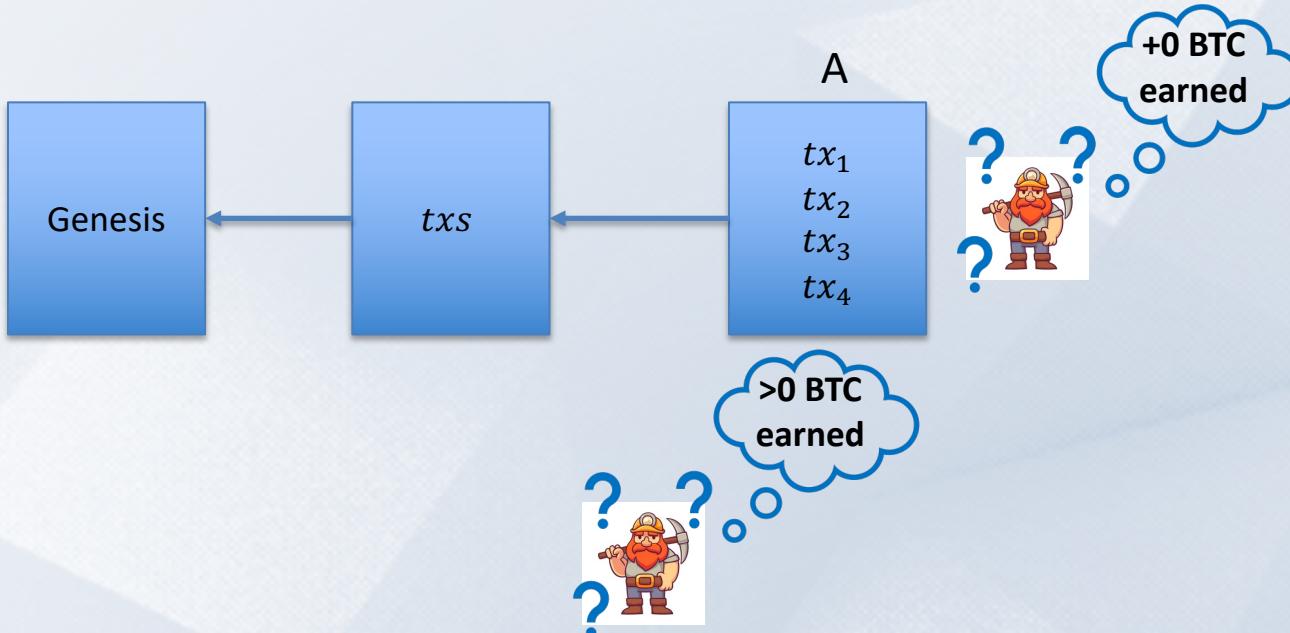
- $e^{-\Omega(k)}$ is the error probability for confirmation.
- Latest result for bounding the error probability as a function of k :

$$\epsilon \leq \left(2 + 2 \sqrt{\frac{1-\beta}{\beta}} \right) (4\beta(1-\beta))^k$$

- We say ‘confirmation’ instead of finalization because when you *confirm* a block or transaction, you *confirm* it with an error probability...
- ...unlike *finalizing* a block where there is no error probability*.

Now, we see why Bitcoin has 1 block every 10 minutes, instead of 1 block every second...

Would $\beta < 1/2$ hold in practice?



Transaction fees
(paid to the miner):

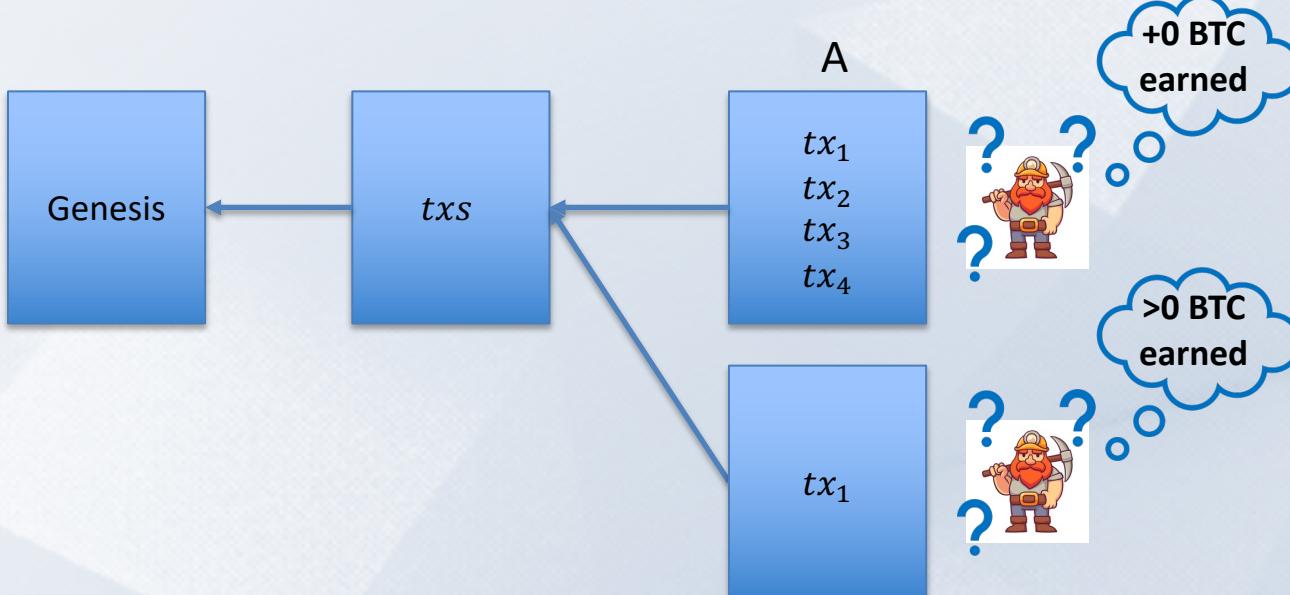
tx_1 : 4 BTC

tx_2 : 3 BTC

tx_3 : 2 BTC

tx_4 : 1 BTC

Would $\beta < 1/2$ hold in practice?



Transaction fees
(paid to the miner):

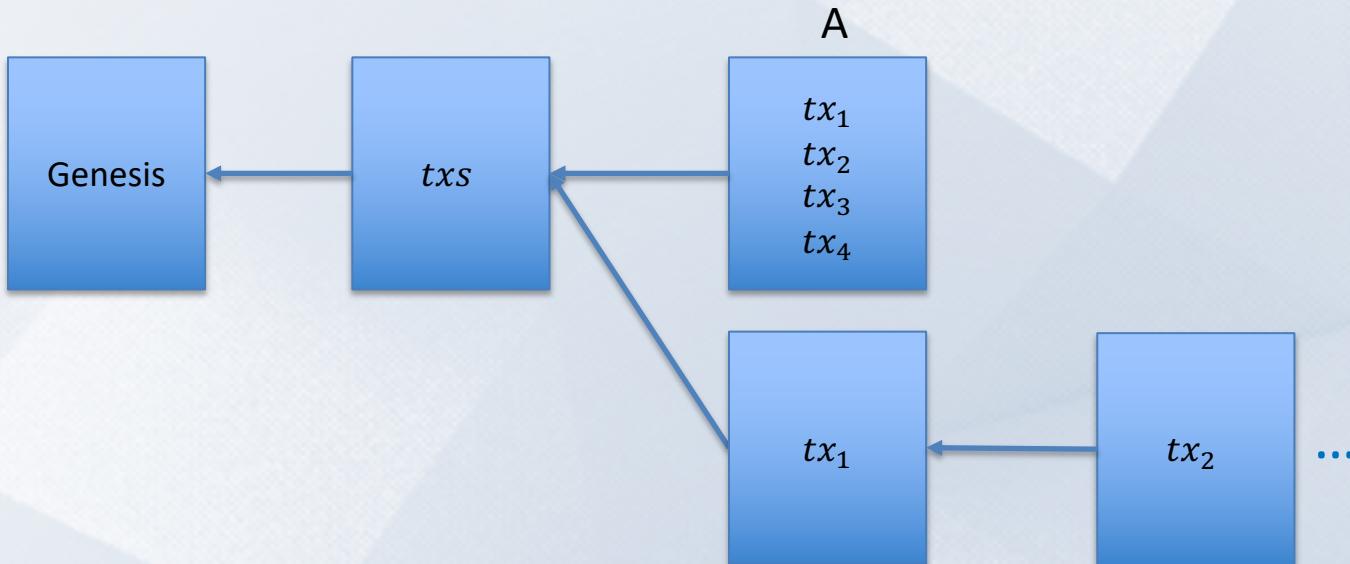
tx_1 : 4 BTC

tx_2 : 3 BTC

tx_3 : 2 BTC

tx_4 : 1 BTC

Would $\beta < 1/2$ hold in practice?



Transaction fees
(paid to the miner):

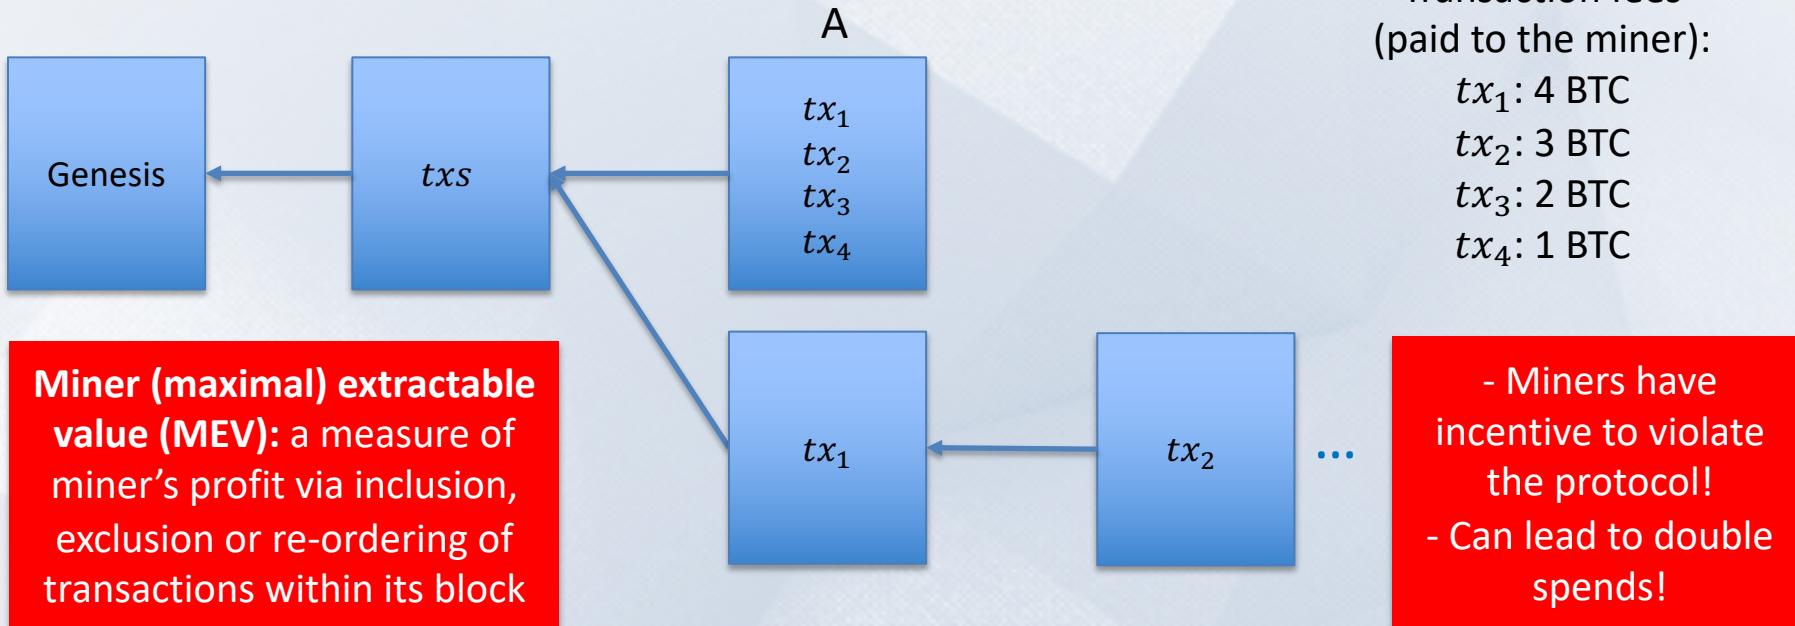
tx_1 : 4 BTC

tx_2 : 3 BTC

tx_3 : 2 BTC

tx_4 : 1 BTC

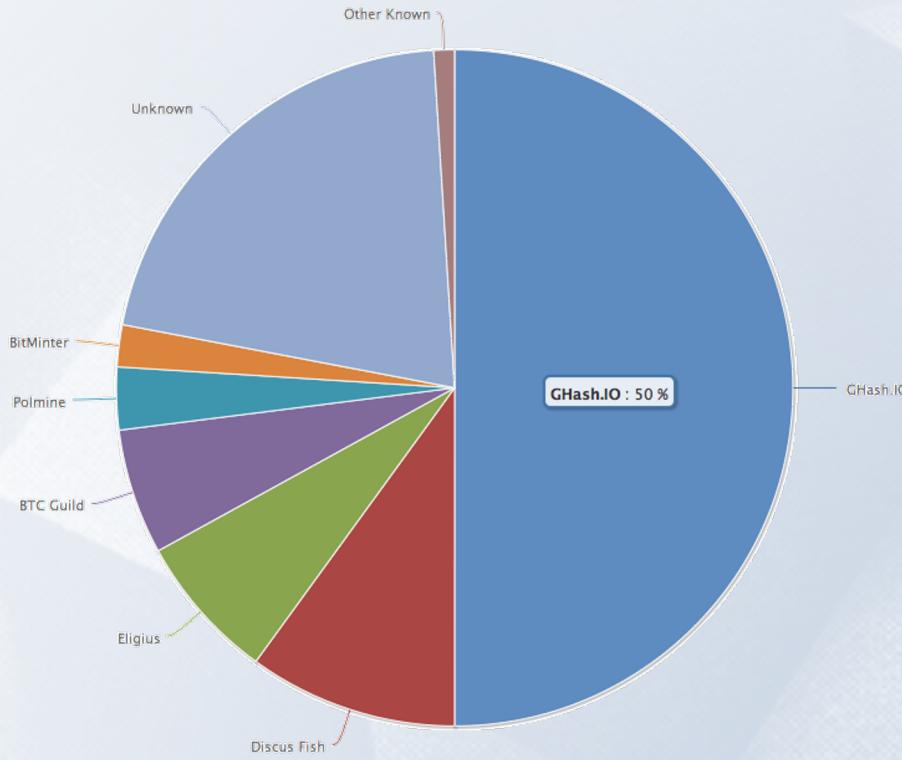
Would $\beta < 1/2$ hold in practice?



Need to think about incentives!!

MEV gives even more incentive to violate the protocol!!

No Attacks on Bitcoin?



Ghash.IO had >50% in 2014

- Gave up mining power

Why are visible attacks not more frequent?

Miners care about the Bitcoin price?

- Not a valid argument.
- They can 'short' the chain for profit!

Might not always be rational to attack.

No guarantees for the future!

Is Bitcoin the Endgame?

Bitcoin provides Sybil resistance and dynamic availability.
Is it the Endgame for consensus?

No!

Bitcoin is secure only under synchrony and loses security during periods of asynchrony.

It confirms blocks with an error probability depending on k , i.e., blocks are not finalized.

Energy consumption?

Next lecture: low-energy consensus using proof-of-stake

END OF LECTURE

Next lecture: Incentives and Accountability in Consensus

